A systematic review of machine learning classification methodologies for modelling passenger mode choice

Tim Hillel [a,b,*], Michel Bierlaire [a], Mohammed Z.E.B. Elshafie [c], Ying Jin [d]

[a] *Department of Engineering, University of Cambridge, UK*
[b] *School of Architecture, Civil and Environmental Engineering, Ecole Polytechnique Fédérale de Lausanne, Switzerland*
[c] *College of Engineering, Qatar University, Doha, Qatar*
[d] *Department of Architecture, University of Cambridge, UK*

ABSTRACT

Machine Learning (ML) approaches are increasingly being investigated as an alternative to Random Utility Models (RUMs) for modelling passenger mode choice. These approaches have the potential to provide valuable insights into choice modelling research questions. However, the research and the methodologies used are fragmented. Whilst systematic reviews on RUMs for mode choice prediction have long existed and the methods have been well scrutinised for mode choice prediction, the same is not true for ML models. To address this need, this paper conducts a systematic review of ML methodologies for modelling passenger mode choice. The review analyses the methodologies employed within each study to (a) establish the state-of-research frameworks for ML mode choice modelling and (b) identify and quantify the prevalence of methodological limitations in previous studies.

A comprehensive search methodology across the three largest online publication databases is used to identify 574 unique records. These are screened for relevance, leaving 70 peer-reviewed articles containing 73 primary studies for data extraction. The studies are reviewed in detail to extract 17 attributes covering five research questions, concerning (i) classification techniques, (ii) datasets, (iii) performance estimation, (iv) hyper-parameter selection, and (v) model analysis.

The review identifies ten common methodological limitations. Five are determined to be methodological pitfalls, which are likely to introduce bias in the estimation of model performance. The remaining five are identified as areas for improvement, which may limit the achieved performance of the models considered. A further six general limitations are identified, which highlight gaps in knowledge for future work.

1. Introduction

Solutions used both in industry and academic research for modelling passenger mode choice have traditionally relied almost exclusively on econometric Discrete Choice Models (DCMs) based on the random utility framework (McFadden, 1981). However, there have been two key recent drivers which have resulted in researchers exploring alternative approaches. Firstly, the adoption of new transportation-related technologies has driven a step change in the availability of data on passenger movements of several orders of

List of acronyms

AB	AdaBoost.
ABM	Agent Based Model.
AI	Artificial Intelligence.
AMPCA	Arithmetic Mean Probability of Correct Assignment.
ANN	Artificial Neural Network.
API	Application Programming Interface.
ASC	Alternative Specific Constant.
BIC	Bayesian Information Criterion.
BL	Bayesian Learner.
BN	Bayesian Network.
CART	Classification and Regression Trees.
CEL	Cross-Entropy Loss.
CNL	Cross-Nested Logit.
DCM	Discrete Choice Model.
DT	Decision Tree.
EL	Ensemble Learning.
ELM	Extreme Learning Machine.
FFNN	Feed-Forward Neural Network.
FL	Fuzzy Logic.
GB	Gradient Boosting.
GBDT	Gradient Boosting Decision Trees.
GPS	Global Positioning System.
IVTT	In-Vehicle Travel Time.
k-NN	k-Nearest Neighbours.
LOS	Level of Service.
LR	Logistic Regression.
ML	Machine Learning.
MLP	Multi-Layer Perceptron.
MSE	Mean Squared Error.
NB	Naive Bayes.
NL	Nested Logit.
O-D	Origin-Destination.
OOB	Out-Of-Bootstrap.
OVTT	Out-of-Vehicle Travel Time.
PNN	Probabilistic Neural Network.
PRISMA	Preferred Reporting Items for Systematic Reviews and Meta-Analyses.
PT	Public Transport.
RBF	Radial Basis Function.
RBFNN	Radial Basis Function Neural Network.
RBML	Rule-Based Machine Learning.
RF	Random Forest.
ROC	Receiver Operating Characteristic.
RP	Revealed Preference.
RSM	Rough Set Model.
RUM	Random Utility Model.
SP	Stated Preference.
SVM	Support Vector Machine.
VOC	Vehicle Operating Cost.
VoT	Value of Time.

magnitude. Secondly, there have recently been significant breakthroughs in Machine Learning (ML) research, which have resulted in numerous success stories of ML applications in other similar tasks.

These drivers have resulted in a number of recent research applications of ML techniques to the mode choice problem. The application of ML has the potential to provide valuable new insights into mode choice modelling research questions. However, the existing research is fragmented, and there have been few studies which comprehensively compare ML techniques with each other and with Random Utility Models (RUMs). Additionally, whilst systematic reviews on RUMs for mode choice prediction have long existed, and the methods have been well scrutinised, the same is not at all true for ML models.

To address these limitations, this paper conducts a systematic review of ML machine learning approaches for modelling passenger mode choice. The paper specifically focuses on classification approaches, where a labelled dataset is used to estimate an individual mode choice model. The review focuses on the methodologies employed within each study, including the classification algorithms, datasets, model validation, model optimisation, and model analysis. The review aims to (a) establish the state-of-research frameworks for ML mode choice modelling and (b) identify and quantify the prevalence of methodological limitations in previous studies. Whilst this review is focused on passenger mode choice literature, the findings are relevant to other choice modelling domains.

2. Machine learning for mode choice prediction

The predominant approach used in industry and academic research for modelling passenger mode choice are Random Utility Models (RUMs) (McFadden, 1981). These models rely on functions of the input variables, called *utility specifications*, for each option (mode) in the choice-set. In a logit model, the utilities (the output values of the utility specifications) are then passed through a logistic function to generate choice probabilities for each option in each considered choice situation. Other model structures, such as the Nested Logit (NL) and Cross-Nested Logit (CNL), allow for these probabilities to be calculated given correlations among the options in the choice-set (Ben-Akiva and Lerman, 1985, Chapter 10). The parameter values in the utility specification are estimated using maximum likelihood estimation, in order to maximise the joint likelihood of the training data given the model.

The utility specifications in the model are defined by the modeller prior to fitting the model. This allows the modeller to incorporate established behavioural theory and expert knowledge into the model. The estimated parameter values can then be used to test hypothesis about the consistency of RUM predictions with expected behaviour. These parameters can also be used to extract key behavioural indicators, such as the elasticities and Value of Time (VoT) (Ben-Akiva and Lerman, 1985, Chapter 5).

The nature of all of the relationships between the input variables and the utilities must be defined in the utility specifications. This includes all non-linear transformations of variables and any interactions between them. As the utility functions are specified in advance of estimating the model, this means that the modeller must hypothesise and test these relationships manually.

In ML terminology, an RUM can be considered as a *supervised probabilistic classifier*; the aim of the model is to predict the probability of an individual choosing each mode (i.e. the *classes*), given a set of *features* (variables) describing the choice situation. The modeller has access to a finite dataset of choice situations alongside the *ground-truth* class labels (the option chosen) to train the model. This task therefore appears to be a natural application for ML classification algorithms, which have shown a great deal of success with similar problems in other research domains, such as image recognition, text classification, and disease detection (Hastie et al., 2008).

Rather than relying on predefined utility specifications, ML classification algorithms instead model the relationship between input features and the class labels directly from the data, without input from the modeller. The majority of ML classifiers (excluding linear models) have the ability to automatically capture non-linear relationships between the inputs and outputs. The added flexibility in ML classifiers compared to RUMs may allow the model to *generalise* to relationships not previously considered and therefore which would not have been identified using manually defined utility specifications.

The greater flexibility of ML classifiers presents a much higher propensity for a model to *overfit* to noise in the training data. Additionally, as there is no underlying behavioural model in an ML model, it is not straightforward to check for or ensure for behavioural consistency of the model predictions, or extract behavioural indicators from the model.

The *generalisation error* measures the ability of a classifier to accurately predict class probabilities for previously unseen data. In ML applications, this is typically estimated by validating model on separate out-of-sample data, unseen by the model during training. The model validation ensures that the model has successfully generalised to valid relationships in the data, without fitting to noise in the data. There is therefore a balance between *underfitting* and *overfitting*, known as the *bias-variance trade-off* (Hastie et al., 2008). If a model has high *bias* it is not flexible enough to identify valid correlations that are present in the real-world test data (underfitting). If a model has high *variance* it is too flexible and is replicating noise in the data without generalising to valid correlations between the input features and class labels (overfitting).

The flexibility of an algorithm to fit to the data when training a model is *regularised* using the algorithm's *hyper-parameters*. These are parameters of the algorithm, such as the maximum permissible number of splits in a decision tree, which impact the bias and variance of the fitted model (see Section 2.1). Model performance is highly dependent on chosen hyper-parameter values, and so it is important to select appropriate values for both the task and data (Hoos et al., 2014).

In order for the model validation to be a true estimate of the generalisation error, the test-set must be truly separate from the training data and not seen by the model at any stage prior to final testing. Incorrect validation approaches can result in *data-leakage*, where the model is somehow exposed to the test-set (potentially including the ground-truth labels) before final testing. This can allow the classifier to fit to the test-set, therefore resulting in the test-error underestimating the generalisation error that would be achieved on truly unseen data. Examples of validation schemes that result in data-leakage include regularising the model based on test performance (e.g. during hyper-parameter selection, see Section 4.4) or through shared information between the train and test-sets (e.g. through inappropriate sampling of hierarchical data, see Section 4.3).

ML classification investigations can be broken into two main processes; *model development* and *model evaluation*. In the model development process, the modeller tries to develop a model with the aim of minimising its generalisation error. This includes hyper-parameter selection, feature processing, and algorithm selection/development. In the model evaluation process, the modeller then estimates the generalisation error of the model, typically by testing on an out-of-sample test-set.

If the model development in a study is not appropriate (e.g. if appropriate hyper-parameter selection is not used) the model will achieve a higher generalisation error than is possible for that algorithm. This means differences in model performance may be due to differences in the model development process, and not to do with the potential performance of the algorithm itself. As such, it is

important to consider the model development process when making comparisons between the relative performance of *algorithms* for a given task. Conversely, if the model evaluation process used is inappropriate (e.g. if there is data-leakage from the test-set during model development) than the estimate of generalisation error will be biased. These issues therefore represent *pitfalls* that will result in unreliable evaluation of model performance. As such, it is important to consider the model evaluation process for any evaluation of *model* or *algorithm* performance for a task.

As there is a lot of overlap in the theory and practice in the fields of RUMs and ML, there are a substantial number of equivalent or nearly-equivalent terms between them. As this paper reviews ML methodologies, the ML terminologies have been preferred. For clarity of the associated material, Table 1 some of the equivalent and nearly equivalent terms used in this paper.

2.1. Machine learning classification algorithms

In order to provide an understanding of the techniques used, the following sections give an overview of five classes of supervised classification algorithm which have previously been used to investigate mode choice, including introducing their main hyper-parameters: Logistic Regression (LR), Artificial Neural Networks (ANNs), Decision Trees (DTs), Ensemble Learning (EL), and Support Vector Machines (SVMs).

Logistic Regression The Logistic Regression (LR) classifier uses the same underlying mathematical formulation as an RUM, with linear functions of the input features passed through the softmax (logistic) function to generate class probabilities. The distinction between the two approaches in that in an RUM regularisation is applied manually through the utility functions, whereas in the ML LR approach only automatic regularisation (or no regularisation) is applied.

Two primary types of regularisation are used in ML LR models. *L1* regularisation (also known as *lasso* regularisation) penalises the model for the sum of absolute values of the weights. Conversely, *L2* regularisation (also known as *ridge* regularisation) penalises the model for the sum of squares of the weights. The amount of regularisation is controlled using the C hyper-parameter, with a larger value of C indicating more regularisation (higher penalty for the values of the weights).

For supervised probabilistic classification using labelled data, intercepts (or Alternative Specific Constants(ASCs) for RUMs) should be included in the model to ensure representative class probabilities. For RUMs, one ASC is typically normalised to zero as an additional constraint to allow for an identifiable solution (Bierlaire et al., 1997). This is not needed in LR models with L1 or L2 regularisation, where the penalty term ensures the solution is uniquely identifiable.

Artificial Neural Networks Artificial Neural Network (ANN) is a term used to cover a family of classifiers which mimic the network structure of the brain. Whilst there are a huge variety of possible ANN structures for dealing with different input data types (e. g. images, time-series, natural language etc), mode choice applications have typically relied on the Feed-Forward Neural Network (FFNN) (also known as the Multi-Layer Perceptron (MLP)) (Svozil et al., 1997).

A FFNN consists multiple *layers* of *nodes* (neurons), including (i) an input layer, which passes the feature values to the network; (ii) an output layer, which outputs the predicted values from the network; and (iii) any number of hidden layers. For probabilistic classification, the number of nodes in the input and output layers is fixed by the number of features and classes in the data respectively. The number of hidden layers and number of nodes in each hidden layer are then hyper-parameters which describe the structure of the network.

Each node has an activation function, which determines the output of that node from the weighted sum of its inputs. This can also be considered as a hyper-parameter. There are many possible activation functions used in practice, including linear, sigmoid, tanh, softplus, softsign, ReLU (rectified linear unit), ELU (exponential linear unit), and SELU (scaled exponential linear unit).

As with RUMs and LR models, the output values of the network are passed through the softmax function to generate classification probabilities. The weights (parameters) for each link in the network are fitted to the input data (equivalent to estimating an RUM).

FFNNs are most commonly trained using *mini-batch gradient descent*. This algorithm splits the input data into small batches. The network weights are then updated iteratively on the individual batches. Each time the model sees all of the data once is termed an *epoch*. The number of epochs can be set to regularise the model and limit overfitting. This hyper-parameter is often set automatically to limit overfitting by applying a stopping criterion based on out-of-sample predictive performance.

In a fully connected network, every node in one layer is linked to every node in the next layer. Further regularisation can be applied using the *dropout* hyper-parameter, which specifies a proportion of the neurons to be dropped randomly from the network for each

Table 1
Equivalent and nearly-equivalent terms between random utility and ML models. ASC = Alternative Specific Constant (ASC).

Random utility	Machine learning	Notes
Attribute	Feature	Variables of the choice-set.
Covariate	Feature	Socio-economic variables of the individual. No distinction is made between attributes and socio-economic covariates in ML classifiers.
ASC	Intercept/bias	Used to ensure representative class proportions for logit models/ANNs estimated on labelled data
Parameter	Weights	Referred to as coefficients in linear utility functions in RUMs. Weights are used only in parametric ML models (LR and ANNs).
Estimate	Train	Both are often referred to as *fitting* the model.
Logistic function	Softmax	Referred to as the sigmoid function in the binary case.

mini-batch of data (Srivastava, Hinton, et al., 2014).

Decision Trees Decision Trees (DTs) (or Classification and Regression Trees (CART)) are classifiers which sort data into groups using a set of sequential splits in a tree-like structure (Breiman, 2017). The most commonly used Decision Trees (DTs) are fitted using recursive binary splits, with each split chosen to result in the greatest reduction in the *randomness* of the data at that point (i.e. it is a *greedy* algorithm). Two metrics can be used to measure how shuffled the data are, *Gini impurity* and *entropy*.

To calculate each split, the data at the selected node are sorted according to each feature, and each possible binary split point (less/greater than a certain value) is tested for each feature. The split point which results in the greatest reduction in the impurity or entropy (across all features) of the data is then selected, resulting in two new child nodes. The same algorithm can then be applied recursively to each child node. This process is repeated until a stopping condition (set using the hyper-parameters) is met. For example, the *maximum depth* specifies the maximum number of sequential splits which can be applied along a branch, the *minimum leaf size* specifies the minimum size *both* nodes of a split must have in order for a split to take place, and the *minimum split size* specifies the minimum number of samples in a node for a split to be considered at that node.

Decision trees can only generate discrete predictions (either classes or finite regression values), and so are not suitable for probabilistic mode choice prediction when used independently. However, they can be combined in ensembles to generate probabilistic predictions.

Ensemble Learning Ensemble Learning (EL) algorithms combine several individual predictive models (called estimators) in an ensemble to improve the quality of predictions. Provided the estimators make errors *independently* (i.e. the learners are uncorrelated), and are more likely to be right than wrong, then combining them in an ensemble reduces their individual uncertainty.

DTs are the predominantly used estimators for (EL). DTs have high variance, making them highly unstable (small changes in the input result in large differences between classifiers). As such, it is relatively easy to train uncorrelated DTs compared to more stable classifiers (e.g. LR). In addition, DTs are algorithmically simple to fit and obtain predictions from. This means that large ensembles of DTs can fit and predict in reasonable time.

Several *meta-algorithms* can be used to combine estimators. This includes algorithms where estimators are trained on the data (or samples of the data) in parallel, e.g. bootstrap aggregating (bagging) and Random Forest (RF), as well as algorithms where the weak learners are estimated sequentially, e.g. AdaBoost (AB) and Gradient Boosting (GB). For ensembles of discrete classifiers, probability-like values can be outputted by calculating the proportions of each class prediction across the estimators in the ensemble. For Gradient Boosting Decision Trees (GBDT), the DTs in the ensemble are trained to output discrete regression values. These values are then summed across the ensemble and passed through the softmax function to output choice probabilities.

As well as the hyper-parameters of the weak learners themselves, the principle hyper-parameter of EL meta-algorithms is the number of estimators in the ensemble. In parallel approaches, this number must be specified. In sequential approaches, a stopping criterion can be applied based on out-of-sample predictive performance (similar to the number of epochs in ANNs).

Support Vector Machine The Support Vector Machine (SVM) algorithm makes use of a *kernel* to transform the data into a high-dimension space. The algorithm then finds the optimal linear decision surface (or *hyper-plane*) in the transformed space which divides the data into two classes (Cortes and Vapnik, 1995).

There are multiple kernels which can be used to transform the data, including linear (no transformation), polynomial, Radial Basis Function (RBF) (or *Gaussian*), and sigmoid.

For linearly-separable data (within the transformed space), the optimal hyperplane is the one that exactly divides the data without misclassification whilst maximising the possible *margin*. The margin is defined as the perpendicular distance between the hyperplane and the nearest data points (these data points are called support *vectors*). For complex, real-world examples, the input data are not normally linearly-separable, even within the transformed space. As such, there is a balance between the width of the hyperplane and the number of misclassifications of the training data. This is controlled using the regularisation parameter (C). A higher value of C represents a higher importance of the misclassified points (higher variance), whilst a lower value of C will put a higher importance on the width of the hyperplane (higher bias).

Support Vector Machines (SVMs) are inherently binary classifiers. However, they can be used for multiclass classification using either a *one-vs-rest* or *one-vs-one* strategy.

SVMs output a continuous score for each prediction. This score can be interpreted as the confidence of the classification. However, these scores do not correspond well to class probabilities (Niculescu-Mizil and Caruana, 2005). Methods to calibrate the scores as class probabilities are proposed by Wu et al. (2004) and Platt (1999).

2.2. Need for a review of machine learning methodologies

As discussed above, ML approaches are increasingly being investigated as an alternative to RUMs for mode choice prediction. However, the research is fragmented, with inconsistent methodologies used in past studies. The implications of different methodological decisions is not yet well understood. As such, there is a need to evaluate the methodologies used in previous studies in order to understand the scope of ML approaches and establish good standard practices.

There exist several review papers in the literature focusing on mode choice modelling, including those by Barff et al. (1982), Hensher and Johnson (1983), Kruger (1991), Nerhagen (2000), Meixell and Norbis (2008), Ratrout et al. (2014), Jing et al. (2018), and Minal and Sekhar (2014). However, all but two of these reviews focus exclusively on statistical RUM techniques. Ratrout, Gazder, and Al-Madani (2014) and Minal and Sekhar (2014) explicitly review ML and Artificial Intelligence (AI) approaches within the literature, including ANN approaches to mode choice modelling alongside RUM based studies. The studies conclude that ANN have been successfully used for mode choice modelling, in particular due to their flexibility when dealing with multidimensional non-linear data.

Ratrout, Gazder, and Al-Madani (2014) further state that whilst the vast majority of existing studies are based on logit models, it can be expected that the trend of using ML methods will continue in future.

Whilst the studies by Ratrout et al. (2014) and Minal and Sekhar (2014) evaluate some of the existing ML mode choice research, they have a number of limitations. Primarily, they focus only on ANN (and Fuzzy Logic (FL)) approaches, and as such do not cover any contributions using other ML techniques, including DTs, SVMs, and EL. Secondly, these reviews are intended to be exploratory as opposed to systematic, and do not represent comprehensive coverage of all relevant studies. Additionally, the reviews are intended to be general, and do not focus on specific aspects of the methodologies used in each study. Finally, there have been a substantial number of new studies published since these reviews were carried out. To address these limitations, this paper conducts a systematic review of ML approaches to passenger mode choice modelling.

2.3. Overview of paper

The remainder of the paper is laid out as follows. Section 3 outlines the methodology for the review, including the research questions, review protocol, and study selection. Next, Section 4 presents the results of the review, first giving an overview of the selected studies, before exploring each research question in turn to identify the methodological limitations. The limitations are categorised into

- *technical limitations:* technical issues within the methodologies of specific studies that are likely to have an impact on their results, which are further categorised into
- *pitfalls:* issues in the model evaluation process which are likely to introduce bias into modelling results, and
- *areas for improvement:* modelling decisions which are not incorrect but could be addressed in order to improve the reliability of the results for comparing the classification algorithms and or the predictive performance of the models; and
- *general limitations:* gaps in knowledge or areas across multiple studies that require further investigative work.

Finally, Section 5 summarises the findings, identifies potential limitations of the review, and presents the conclusions.

3. Methodology

The procedure for this systematic review is adapted from that given by Kitchenham and Charters (2007). The suggested procedure suggested has 10 stages broken down into three phases:

- *Planning the review*
  1. Identification of the need for a review
  2. Specifying the research questions
  3. Developing a review protocol
- *Conducting the review*
  4. Identification of research
  5. Selection of primary studies
  6. Study quality assessment
  7. Data extraction and monitoring
  8. Data synthesis
- *Reporting the review*
  9. Specifying dissemination mechanisms
  10. Formatting the main report

This review is focused on summarising the methodologies used in each study, and as such, no attempt is made to draw conclusions from the aggregate results or combined findings of the studies. Consequently, no assessment of the quality of each study is made (step 6 in the framework). The review presented in this paper therefore consists of the nine remaining stages presented above.

The focus of this review is the methodologies used in ML approaches to modelling passenger mode choice. In particular, the review serves to investigate the following research questions:

1. Which classification techniques have been used to investigate mode choice?
2. What is the nature of datasets used to investigate mode choice?
3. How is model performance determined?
4. How are optimal model hyper-parameters selected?
5. How is the best model selected?

3.1. Review protocol

This section outlines the protocol for the search strategy, selection criteria, and data extraction strategy.

3.1.1. Search strategy

The search strategy is used to identify relevant papers to the review. In order to ensure full coverage of relevant papers, papers are collated from three databases: the two major online curated publication databases, Web of Science and Scopus; and the Google Scholar search engine. The same search is repeated for each database.

In order to only select papers that discuss ML techniques, only papers with one or more selected phrases relating to ML across all relevant fields are selected. The following initial phrases are tested: *machine learning*, *neural network*, *decision tree*, *ensemble method*, *random forest*, *boosting*, and support *vector*.

This review focuses on papers with a core focus of mode choice modelling. As such, only papers with the *title* directly relating to mode choice are included. The following initial phrases are tested: *mode choice*, *mode selection*, *travel mode*, *transport mode*, *transportation mode*, and *mode of travel*. The requirement of having one of the mode choice phrases in the title is used to automatically pre-screen irrelevant papers. A Google scholar search for papers containing at least one of the above ML phrases alongside at least one of the mode choice phrases across all relevant fields returns over 15000 results (as of December 2019).

Papers from any period up until the search date are included in the search.

The initial search phrases are tested in different combinations across the three databases. The terms *mode of travel* and *mode selection* are omitted from the title search, as they return no relevant papers when used alongside the ML search terms.

Additionally, a number of papers using Fuzzy Logic (FL) (within Rule-Based Machine Learning (RBML)) were found in the initial search results. To reflect this, the phrase *fuzzy logic* is added to the search across all relevant fields.

3.1.2. Selection criteria

The following eligibility criteria are determined for the papers found in the search to be included in the study:

- Studies in peer-reviewed journals or conference proceedings written in English
- Studies which investigate passenger mode choice at disaggregate (individual) level.
- Studies which employ one or more ML technique(s) for predictive modelling.

Paper selection is carried out using the Preferred Reporting Items for Systematic Reviews and Meta-Analyses (PRISMA) guidelines (Moher et al., 2009). Firstly, duplicates are removed from the search records. Secondly, the record titles and abstracts are screened against the eligibility criteria. Finally, the remaining full-text articles are assessed for eligibility. All stages of the selection criteria are carried out independently by the first author.Where a paper contains more than one relevant modelling scenario (defined as having separate input datasets and different methodologies), each modelling scenario is treated as a separate study for the analysis.

3.1.3. Data extraction strategy

In order to extract the necessary data from each study without bias, a list of attributes is collected from each study. The attributes, shown in Table 2, are intended to be specific, objective, and quantifiable/categorical, in order to limit subjectivity in the data

Table 2
Research questions and corresponding attributes of studies for data extraction.

No.	Description
Q1	**Which classification techniques have been used to investigate mode choice?**
Q1a	Classification algorithms used in study
Q1b	Logit model implementation
Q2	**What is the nature of datasets used to investigate mode choice?**
Q2a	Nature of dataset
Q2b	Unit of analysis
Q2c	Dataset availability
Q2d	Modes in choice-set
Q2e	Modelling of mode-alternatives
Q2f	Input features dependent on output choice
Q2g	Hierarchical data
Q3	**How is model performance determined?**
Q3a	Validation method
Q3b	Sampling method
Q3c	Performance metrics used
Q4	**How are optimal model hyper-parameters selected?**
Q4a	Hyper-parameter search method
Q4b	Hyper-parameter validation method
Q4c	Hyper-parameter validation data
Q5	**How are the final models analysed?**
Q5a	Statistical testing
Q5b	Extraction of behavioural indicators

extraction process. Together the attributes provide the evidence for the research questions.

Data extraction is carried out independently by the authors. Each study is reviewed in detail, with each attribute for each study determined and tabulated in a spreadsheet. Separate entries are entered into the spreadsheet for papers containing multiple studies (modelling scenarios).
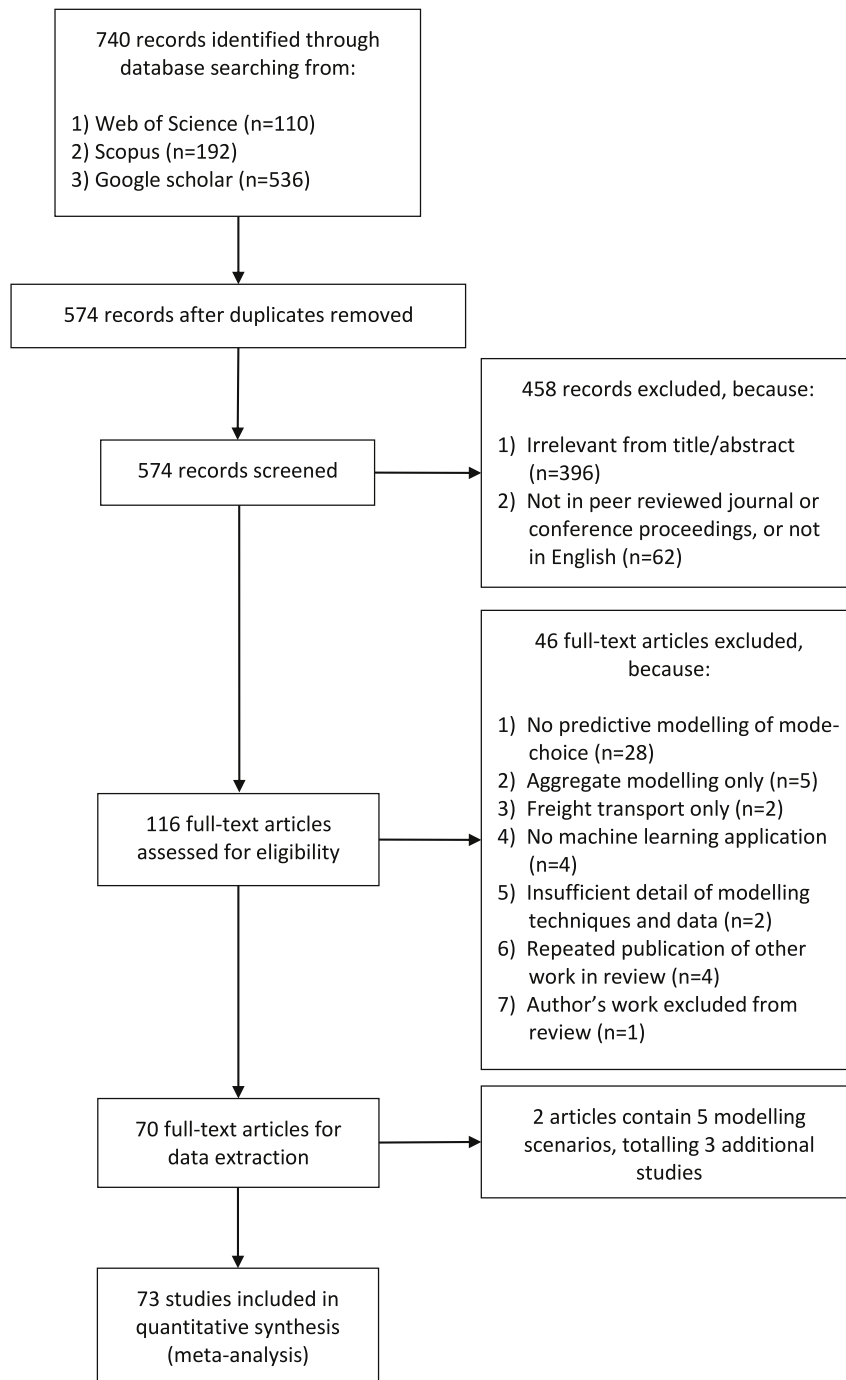

Fig. 1. PRISMA flowchart of study selection process.

3.2. Study selection

The following search terms are used to carry out the search strategy outlined in Section 3.1.

- **Web of Science**: *TITLE: ("mode choice" OR "travel mode" OR "transport mode" OR "transportation mode") AND TOPIC: ("machine learning" OR "neural network" OR "decision tree" OR "ensemble method" OR "random forest" OR "boosting" OR "support vector" OR "fuzzy logic")*
- **Scopus**: *(TITLE ("mode choice" OR "travel mode" OR "transport mode" OR "transportation mode") AND TITLE-ABS-KEY ("machine learning" OR "neural network" OR "decision tree" OR "ensemble method" OR "random forest" OR "boosting" OR "support vector" OR "fuzzy logic"))*
- **Google Scholar**: (intitle:"mode choice" OR intitle:"travel mode" OR intitle:"transport mode" OR intitle:"transportation mode") AND ("machine learning" OR "neural network" OR "decision tree" OR "ensemble method" OR "random forest" OR "boosting" OR "support vector" OR "fuzzy logic").

  Due to the restriction on search length in Google Scholar, this search is divided into two separate searches, with the results combined.

The search was carried out on 20/12/2019 on all three databases. Fig. 1 shows a PRISMA flowchart of the study selection process. There were 110 records returned from the Web of Science search, 192 records from Scopus, and 536 records from Google Scholar, for a total of 838 records. Duplicates are then removed, leaving 574 records to be screened. The total number of records after removing duplicates is more than were obtained from any one database, showing that there were results from Web of Science/Scopus which were not returned with the Google Scholar search.

The 574 remaining records are then screened as to whether they meet the eligibility criteria outlined in Section 3.1. During screening, 396 papers are excluded for relevance on the basis of their title and abstract. The majority of these records relate to transportation mode detection from Global Positioning System (GPS) data. Of the records which are deemed relevant, a further 62 are excluded as they are not published in peer reviewed publications, (e.g. Thesis/dissertation, unpublished paper, book section), or are not written in English (only having a title and abstract in English).

Table 3
Selected primary articles for review.

No.	Paper	No.	Paper
S1	Raju et al. (1996)	S36	Ermagun and Ansari Lari (2015)
S2	Subba Rao et al. (1998)	S37	Gazder and Ratrout (2015)
S3	Hensher and Ton (2000)	S38	Jia et al. (2015)
S4	Cantarella and de Luca (2003)	S39	Kedia and Krishna (2015)
S5	Van Middelkoop, Borgers, and Timmermans (2003)	S40	Ma (2015)
S6	Xie et al. (2003)	S41	Omrani (2015)
S7	Karlaftis (2004)	S42	Papaioannou and Martinez (2015)
S8	Cantarella and de Luca (2005)	S43	Pitombo and Souza (2015)
S9	Andrade et al. (2006)	S44	Tang et al. (2015)
S10	Shafahi and Nazari (2006)	S45	Li et al. (2016)
S11	Edara et al. (2007)	S46	Sekhar et al. (2016)
S12	Errampalli et al. (2007)	S47	Semanjski et al. (2016)
S13	Moons et al. (2007)	S48	Hagenauer and Helbich (2017)
S14	Wang and Moon (2007)	S49	Hussain et al. (2017)
S15	Zhang and Xie (2008)	S50	Juremalani (2017)
S16	Biagioni et al. (2009)	S51	Lindner and Souza Pitombo (2017)
S17	Chalumuri et al. (2009)	S52	Ma et al. (2017)
S18	Seetharaman et al. (2009)	S53	Nam et al. (2017)
S19	Lu and Kawamura (2010)	S54	Assi and Nahiduzzaman, 2018
S20	Zhao et al. (2010)	S55	Ding et al. (2018)
S21	Xian-Yu (2011)	S56	Golshani et al. (2018)
S22	Yin and Guan (2011)	S57	Lee et al. (2018)
S23	Zenina and Borisov (2011)	S58	Liang et al. (2018)
S24	Zhou and Lu (2011)	S59	Srivastava and Ravi Sekhar (2018)
S25	Dell'Orco and Ottomanelli (2012)	S60	Wang and Ross (2018)
S26	Tang et al. (2012)	S61	Zhu et al. (2018)
S27	Gao et al. (2013)	S62	Assi et al. (2019)
S28	Kumar et al. (2013)	S63	Chang et al. (2019)
S29	Omrani et al. (2013)	S64	Chapleau et al. (2019)
S30	Pulugurta et al. (2013)	S65	Cheng et al. (2019)
S31	Ramanuj and Gundaliya (2013)	S66	Minal et al. (2019)
S32	Shukla et al. (2013)	S67	Pirra and Diana (2019)
S33	Cheng et al. (2014)	S68	Wang and Zhao (2019)
S34	Hossein et al. (2014)	S69	Yang and Ma (2019)
S35	Rasouli and Timmermans (2014)	S70	Zhou et al. (2019)

The full text is obtained for the remaining 116 articles for further review. Of these, a further 46 are excluded on the basis of the selection criteria, as detailed in Fig. 1. This leaves 70 selected articles for data-extraction.

Two articles contain multiple modelling scenarios, for a total of 73 separate studies for meta-analysis.

4. Results and discussion

This section presents the results obtained from the systematic review process. Firstly, sec:articles provides an overview of the 70 articles used for data extraction, including the publication sources and years. The articles with multiple studies are identified, and each of the 73 studies are given a unique identifier. Sections 4.1 to 4.4 then use evidence from the 73 studies to explore each of the five research questions in turn.

4.1. Articles for data extraction

This section provides an overview of the 70 articles used for data extraction. Table 3 provides a unique identifier for each article, alongside its individual reference.

Two papers [S7; S8] contain multiple modelling scenarios, using separate datasets and a different methodology for each one. A separate identifier is assigned to each modelling scenario in each of these papers, and they are treated as separate studies for the meta-analysis. Table 4 provides the label for the additional studies, alongside a description of each modelling scenario. The two papers have a total of five modelling scenarios. This results in a total of 73 studies for meta-analysis.

Five further papers have multiple modelling phases but are deemed not to be separate studies for the purpose of this review.

S3, S17 and S63 each include input datasets for two separate cities. In each of these papers, the datasets are collected and analysed using very similar methodologies, and so are treated as part of the same study for the purpose of this review. The largest dataset in each paper is used for the analysis in Section 4.2: the combined dataset of Sydney and Melbourne for S3, the Visakahpatnam dataset for S17, and the German nationwide dataset for S63.

S34 and S37 each include three separate modelling phases. In S34 each phase represents a choice in a sequence for tour-based mode choice. *Model 2-1*, which predicts attributes of the first trip in a day made by an individual, is analysed within this review. In S37 each phase models different choice situations using the same modelling methodology on subsets of the same dataset. *Phase I*, which models the revealed preference choice between car and plane, is used for the analysis in the review.

4.1.1. Publication source

Table 5 provides details of all journals and conferences/proceedings from which more than one article was selected. The articles come from a wide spread of publications, with a total of 33 different journals and 16 different conferences featured. The majority of the papers (45/70) are published in journals, making up 64% of the articles, with the remaining 25 papers (36%) published in conference proceedings.

The top two sources for articles are the Transportation Research Record Journal and the Transportation Research Board Annual Meeting conference, both of which are published by the Transportation Research Board. Together, they make up 20% (14/70) of the articles.

4.1.2. Publication year

Fig. 2 shows the distribution of article publication dates from 1995 to 2019.

There is a clear upwards trend of increasing number of publications regarding ML applications to mode choice per year. Half of the selected articles were published from 2015 onwards. Conversely, only 10 relevant papers were published prior to 2007.

4.2. Which classification techniques have been used to investigate mode choice?

The following sections present an overview of the classification techniques used in the 73 studies in the review.

Q1a: Classification algorithms used in study

Based on the responses to Q1a, the classification techniques are grouped into nine categories, as shown in Table 6. A brief overview of the classification techniques identified in this paper is given in Section 2.1. For each algorithm, an example paper from the

Table 4
Primary studies with multiple modelling scenarios in review.

No.	Paper	No.	Scenario
S7	Karlaftis (2004)	S7.1	Interurban mode choice in Australia
		S7.2	Commuter mode choice in Athens, Greece
		S7.3	Commuter mode choice in Las Condes-CBD corridor, Chile
S8	Cantarella and de Luca (2005)	S8.1	VENETO dataset
		S8.2	UNISA dataset

Table 5

Summary of publication sources contributing more than one paper to review. Multi-conference proceedings are shown in bold, with the individual conferences in italics below.

Publication	Type	No.
Transportation Research Record	Journal	8
Transportation Research Board Annual Meeting	Conference	7
Transportation Research Procedia:	**Proceedings**	**4**
Euro Working Group on Transportation	*Conference*	*3*
Transportation Planning and Implementation Methodologies for Developing Countries	*Conference*	*1*
Travel Behaviour and Society	Journal	3
International Journal for Traffic and Transport Engineering	Journal	2
Transportation Planning and Technology	Journal	2
Transportmetrica A: Transport Science	Journal	2
East Asia Society for Transportation Studies	Conference	2
International Conference of Chinese Transportation Professionals	Conference	2
Totals (all papers)	**Journal**	46
	Conference	26

systematic review which makes use of that algorithm is provided.

Table 7 shows which classification techniques are used in each study. The majority of studies (47/73) compare ML techniques with statistical RUMs and LR, making logit models the most commonly used classification technique in the studies. The most commonly used ML algorithms are ANNs (34 studies).

Table 7 also shows an increasing focus in the literature on EL (7/15 studies published in the last two years) and SVMs (5/15 studies published the last two years).

Q1b: Logit model implementation

Whilst the overall focus of this review is the ML methodologies used in the studies, Q1a identifies 47 studies which compare ML approaches with logit models (statistical RUMs and (LR)). As such, this section gives a brief overview of the logit models used in these studies.

Table 8 summarises the regularisation method used for the logit model or models in each study, as well as whether intercepts or ASCs are included in the model specification.

As discussed in Section 2.1, a distinction is made between RUMs, where the model is regularised manually through the use of utility specifications and LR, where either no regularisation or L1/L2 regularisation (or a combination of the two) is used. A logit model is classed as using manual utility specification regularisation if significance testing is used to remove any variables or if the utility functions used for each mode are not uniform (including the use of alternative specific Level of Service (LOS) variables for each mode).

Fig. 2. Publication distribution of articles in systematic review (a) per year and (b) cumulative.

Of the 47 studies which use logit models, 22 regularise the model through the use of manually specified utility functions [S2; S3; S4; S8.1; S8.2; S9; S15; S17; S18; S21; S30; S36; S37; S43; S44; S45; S53; S54; S56; S57; S59; S60]. A single study [S70] makes use of an LR classifier with ML regularisation, though it is not stated whether L1 or L2 regularisation is used. 11 studies [S6; S13; S24; S40; S48; S49; S51; S52; S55; S58; S65] make use of no regularisation and include all variables uniformly for all modes. The remaining 13 studies [S7.2; S12; S16; S23; S25; S29; S33; S41; S50; S61; S64; S66; S68] do not describe the logit modelling in sufficient detail to ascertain either regularisation was applied to the model (either through utility functions or L1/L2 regularisation).

Regarding the use of ASCs/intercepts, one study [S2] includes an RUM specification with no ASCs. A further 19 studies [S7.2; S12; S13; S16; S23; S25; S29; S33; S37; S41; S48; S50; S55; S61; S64; S65; S66; S68; S70] do not describe the model in enough detail to ascertain whether intercepts/ASCs are included in the model.

4.2.1. Techniques - limitations

One general limitation is identified regarding the ML techniques used to investigate mode choice: the inconsistent representation of logit models in ML studies. Q1b highlights the inconsistent representation of logit models in the studies in the review. Twenty-four studies either use no regularisation for the logit model, or do not provide sufficient information to gather whether regularisation is used. A further 20 models either do not include intercepts or ASCs in the utility specifications, or do not describe the model in enough detail to verify whether they are included. This is despite intercepts/ASCs being a necessity to reproduce representative choice probabilities for labelled data. In order to make valid comparisons between RUMs and LR with other ML classifiers, it is essential a valid model specification is used.

4.3. What is the nature of datasets used to investigate mode choice?

The following sections discuss the datasets used in the 73 studies in the review, focusing in turn on the nature of the dataset (trip

Table 6
Classification techniques used in studies in review.

Classification algorithm	Example reference
1. Logit models (Log)	
Logistic Regression (LR)	Cantarella and de Luca (2005)
Nested Logit (NL)	Hensher and Ton (2000)
Cross-Nested Logit (CNL)	Nam et al. (2017)
2. Artificial Neural Networks (ANNs)	
Feed-Forward Neural Network (FFNN)	Lee et al. (2018)
Radial Basis Function Neural Network (RBFNN)	Omrani (2015)
Probabilistic Neural Network (PNN)	Zhou and Lu (2011)
Extreme Learning Machine (ELM)	Assi et al. (2019)
Other neural network structures	Cantarella and de Luca (2003)
3. Decision Trees (DTs)	Karlaftis (2004)
4. Extreme Learning Machine (ELM)	
Random Forests (RFs)	Rashidi et al. (2014)
Gradient Boosting (GB)	Wang and Ross (2018)
AdaBoost (AB)	Biagioni et al. (2009)
Bagging	Hagenauer and Helbich (2017)
5. Support Vector Machines (SVMs)	Xian-Yu (2011)
6. Bayesian Learners (BLs)	
Naïve Bayes (NB)	Hagenauer and Helbich (2017)
Bayesian Network (BN)	Ma (2015)
Tree Augmented Naïve Bayes	Tang et al. (2012)
7. Rule-Based Machine Learning (RBML)	
Fuzzy Inference System	Dell'Orco and Ottomanelli (2012)
Rough Set Model (RSM)	Cheng et al. (2014)
Class Association Rules	Lu and Kawamura (2010)
8. Hybrid methods (HM)	
Clustered Logistic Regression	Li et al. (2016)
Logit-ANN	Gazder and Ratrout (2015)
Mixed classifier ensembles	Chang et al. (2019)
9. Miscellaneous (Msc)	
Multivariable Fractional Polynomials	Nam et al. (2017)
Discriminant Analysis	Karlaftis (2004)
Structural Equation Modelling	Papaioannou and Martinez (2015)
Linear regression	Ramanuj and Gundaliya (2013)
k-Nearest Neighbours (k-NN)	Zhou et al. (2019)

Table 7
ML techniques used in each study in review.

No.	Log	ANN	DT	EL	SVM	BL	RBML	HM	Msc	No.	Log	ANN	DT	EL	SVM	BL	RBML	HM	Msc
S1		✓								S35				✓					
S2	✓	✓								S36	✓			✓					
S3	✓	✓								S37	✓	✓						✓	
S4	✓	✓								S38		✓			✓				
S5			✓							S39							✓		
S6	✓	✓	✓							S40	✓					✓			
S7.1			✓							S41	✓	✓			✓				
S7.2	✓	✓	✓						✓	S42									✓
S7.3			✓							S43	✓			✓					
S8.1	✓	✓								S44	✓			✓					
S8.2	✓	✓								S45	✓							✓	
S9	✓						✓			S46		✓	✓						
S10							✓			S47					✓				
S11		✓								S48	✓	✓			✓	✓	✓		
S12	✓						✓			S49	✓	✓							
S13	✓		✓		✓				✓	S50	✓				✓	✓			
S14							✓			S51	✓	✓	✓						
S15	✓	✓			✓					S52	✓						✓		
S16	✓		✓	✓	✓	✓				S53	✓	✓							
S17	✓	✓								S54	✓	✓							
S18	✓						✓			S55	✓				✓				
S19							✓			S56	✓	✓							
S20		✓								S57	✓	✓							
S21	✓	✓								S58	✓				✓				
S22		✓								S59	✓	✓							
S23	✓		✓						✓	S60	✓				✓				
S24	✓	✓								S61	✓			✓		✓			
S25	✓						✓			S62		✓			✓			✓	
S26		✓				✓				S63			✓	✓	✓	✓		✓	
S27		✓								S64	✓			✓					
S28							✓			S65	✓				✓	✓			
S29	✓	✓	✓		✓	✓			✓	S66	✓	✓					✓	✓	
S30	✓						✓			S67					✓				
S31		✓							✓	S68	✓	✓							
S32		✓	✓							S69					✓				
S33	✓						✓			S70	✓		✓	✓	✓	✓			✓
S34			✓	✓						Sum	47	34	18	14	15	9	12	5	7

diary/single-trip questionnaire/stated preference survey, etc); the unit of analysis (trip/tour/commute pattern/mobility); the size of the dataset; the dataset availability; the modes in the choice-set; the modelling of mode-alternatives; input features dependent on output choice; and hierarchical data.

Q2a: Nature of dataset

Table 9 shows the description and size of each dataset.

Only four studies [S3; S9; S53; S68] use (SP) data. One study [S18] uses synthetic choice data, where the choice for a hypothetical metro service is synthesised based on a proposed fare structure and the respondent's willingness-to-pay (which is recorded during the interview).

The remaining 68 studies use Revealed Preference (RP) data. One study [S70] makes use of Origin-Destination (O-D) pairs collected from taxi GPS and bike-sharing scheme data. The remaining 67 studies use datasets specifically collected to investigate mode choice, either from trip-diaries or or single-trip questionnaires.

Thirty-six studies make use of trip diary or activity-diary data, over periods ranging from one day to one year. These diaries are collected either from household surveys [S6; S11; S16; S19; S30; S32; S33; S34; S35; S39; S43; S44; S48; S51; S55; S56; S57; S58; S60; S61; S63; S64; S65; S67; S69] or individual surveys [S5; S13; S40; S42; S47; S52]. Five studies which use trip diary data do not specify enough detail to determine if an individual or household survey is used [S15; S21; S24; S26; S27].

In many studies, a subset of trips is selected from complete trip diaries, e.g. work trips only [S1; S6; S13; S15; S19; S21; S52; S55], education trips only [S39], shopping/social trips only [S56; S61], outbound trips only [S33], trips from home only [S57], first trip of the day only [S34], morning peak trips only [S64], or random sampling [S24; S26; S27].

Twenty studies use individual single-trip questionnaires, where an individual is asked about a single trip they have made [S2; S4; S7.1; S7.2; S7.3; S8.1; S8.2; S14; S20; S23; S25; S36; S37; S45; S49; S59] or a commute they make regularly [S28; S50; S54, S62].

Table 8
Summary of logit models used in each study in review. Reg. = Regularisation, Int. = Intercept included, UF=Utility functions.

No.	Reg.	Int.	No.	Reg.	Int.	No.	Reg.	Int.
S2	UF	✓	**S24**	×	✓	**S52**	×	✓
S3	UF	✓	**S25**	?	?	**S53**	UF	✓
S4	UF	✓	**S29**	?	?	**S54**	?	?
S6	×	✓	**S30**	UF	✓	**S55**	UF	✓
S7.2	?	?	**S33**	?	?	**S56**	×	?
S8.1	UF	✓	**S36**	UF	✓	**S57**	UF	✓
S8.2	UF	✓	**S37**	UF	?	**S58**	UF	✓
S9	UF	✓	**S40**	×	✓	**S59**	×	✓
S12	?	?	**S41**	?	?	**S60**	UF	✓
S13	×	?	**S43**	UF	✓	**S61**	UF	✓
S15	UF	✓	**S44**	UF	✓	**S64**	?	?
S16	?	?	**S45**	UF	✓	**S65**	×	?
S17	UF	✓	**S48**	×	?	**S66**	?	?
S18	UF	✓	**S49**	×	✓	**S68**	?	?
S21	UF	✓	**S50**	?	?	**S70**	L1/L2	?
S23	?	?	**S51**	×	✓			

Two studies [S29; S41] make use of a household survey, in which each working member of the household details their work commute.

Nine studies [S1; S10; S12; S17; S22; S31; S38; S46; S66] do not describe the data in enough detail to be able to determine the nature of the dataset.

The size of each dataset is also shown in Table 9. One study [S59] uses a dataset with under 100 entries. Twenty studies use small datasets, with between 100 and 1000 entries. Thirty-four studies use medium datasets, with between 1000-10 000 entries. Ten studies use large datasets, with between 10000-100 000 entries. Six studies use datasets larger than 100 000 entries.

Two studies [S30; S35] do not give the exact size of the dataset. They both use the trip diaries of individuals in a household survey (5822 individuals in S30, 1446 individuals in S35).

Q2b: Unit of analysis

Sixty-seven of the studies use a single independent choice as the unit of analysis. The choice can be for a single one-way trip per respondent, a return trip (by assuming each leg is made by the same mode), trip diary data where sequences of trips are treated as independent, a regular commute, or a stated preference. Sixty-five of these studies model the mode choice only, whilst two studies [S34; S61] jointly consider other trip attributes (see Q2e).

Six studies use a different unit of analysis. Four studies analyse *mobility*. S51 and S58 both analyse household mobility by predicting the predominant mode used by a household across all trips made on the survey day. S43 analyses individual mobility, by predicting the predominant mode used by an individual across all trips they make on the survey day. Finally, S11 analyses the mobility within clusters. Clusters of similar trips are generated using *k*-means clustering (Hartigan and Wong, 1979). The proportions of trips made by each mode within these clusters is then predicted.

Two studies use a tour-based approach. S16 uses the predicted mode choice of the first trip in a tour (the *anchor mode*) as an input feature for subsequent trips. S67 groups trips into home-based tours across eight categories and predicts overall mode choice for each tour (including mixed mode tours).

Note that (as discussed in Section 4) S34 implements a tour-based analysis, but the subsequent trips in a tour are predicted on the basis of the attributes of the previous trip (including mode choice) as recorded in the dataset, and not as predicted by a model. As such, only the model which predicts attributes of the first trip of the day is analysed in the review (*Model 2-1* in the paper).

Q2c: Dataset availability

An attempt was made to identify and check the availability of the dataset used in each study. The following section discusses all datasets which were found to be openly available. Note that some studies which make use of open data may not have been identified, due to resource constraints when searching for datasets (see Section 5).

Twenty-one studies are identified as using open or partially open data. The majority (11) use openly available household travel

Table 9

Nature and size of dataset used in each study in review.

No.	Type	N
S1	Unclear household survey (work trips only, 535 trips sampled randomly from 3 500)	535
S2	Individual single-trip questionnaire (access to rail on work trip)	4335
S3	Stated preference - individual panel survey (3 trips per person)	801
S4	Individual single-trip questionnaire (mixed purpose urban)	2350
S5	Trip diaries from individual survey (1 year, >2 day vacations only, 7 121 vacations, 2 791 individuals)	7121
S6	Trip diaries from household survey (2-day, 4 746 outbound work trips)	4746
S7.1	Individual single-trip questionnaire (mixed-purpose)	210
S7.2	Individual single-trip questionnaire (mixed purpose urban)	7100
S7.3	Individual single-trip questionnaire (morning home-work trip)	617
S8.1	Individual single-trip questionnaire (student extra-urban trips)	1116
S8.2	Individual single-trip questionnaire (mixed purpose urban)	2350
S9	Stated preference - individual panel survey (160 individuals, 6 trips per individual)	960
S10	Unclear household survey	4147
S11	Trip diaries from household survey (1 year, >100 mile, business trips only)	118000
S12	Unclear individual survey	2868
S13	Trip (Activity) diaries from individual survey (commute patterns extracted)	1025
S14	Individual single-trip questionnaire (fixed O-D, mixed-purpose)	366
S15	Trip diaries from unclear survey (outbound work trip only)	5029
S16	Trip diaries from household survey (1–2 day, 116 666 trips, 19 118 tours)	116666
S17	Unclear survey	1045
S18	Individual single-trip questionnaire (synthetic choice)	229
S19	Trip diaries from household survey (1-day, morning-peak home-work trips only)	9210
S20	Individual single-trip questionnaire	100
S21	Trip diaries from unclear survey (work travel mode choice)	4725
S22	Unclear survey	1007
S23	Individual single-trip questionnaire (mixed-purpose)	498
S24	Trip diaries from unclear survey (500 trips sampled from larger survey, 125 for each mode)	500
S25	Individual single-trip questionnaire (outbound home-work trip)	361
S26	Unclear daily travel survey (2000 trips sampled from larger survey, 500 for each mode)	2000
S27	Trip diaries from unclear survey (650 trips sampled from larger survey, 130 for each mode)	650
S28	Individual single-trip questionnaire (work commute)	606
S29	Commute patterns in household economic survey (9 500 individuals, 3 670 households)	3673
S30	Trip diaries from household survey (Unknown trips, 5 822 individuals, 2 627 households)	?
S31	Unclear household survey	1348
S32	Trip diaries from household survey (1-day)	100000
S33	Trip diaries from household survey (5 721 outbound trips only, 4 831 individuals, 1809 households)	5721
S34	Trip diaries from household survey (1 day, only first trip, 24 807 individuals, 12 568 households)	24807
S35	Trip diaries from household survey (1-day, unknown trips, 1 446 individuals)	?
S36	Individual single-trip questionnaire (outbound home-school trip)	4700
S37	Individual single-trip questionnaire (cross-border)	516
S38	Unclear trip survey (4 500 trips sampled from 17 539)	4500
S39	Trip diaries from household survey (education trips only)	409
S40	Trip diaries from individual survey (1-day, 11 993 trips made by 7 235 people)	11993
S41	Commute patterns in household economic survey (9 500 individuals, 3 670 households)	3670
S42	Trip diaries from individual survey (530 trips, <382 individuals)	530
S43	Trip diaries from household survey (1-day, mobility of household head only)	1216
S44	Trip diaries from household survey (2-day, 72 536 trips, ~31 000 individuals, ~14 000 households)	72536
S45	Individual single-trip questionnaire (holiday travel)	731
S46	Unclear household survey	5843
S47	Trip diaries from individual GPS survey (4 months, 17 040 trips, 292 individuals)	17040
S48	Trip diaries from household survey (6 day, 230 608 trips, 69 918 individuals)	230608
S49	Individual single-trip questionnaire (mixed purpose urban)	620
S50	Individual single-trip questionnaire (work commute)	224
S51	Trip diaries from household survey (mode choice analysed at household mobility level)	18733
S52	Trip diaries from individual survey (1-day, commute patterns extracted)	5040
S53	Stated preference - individual panel survey	6768
S54	Trip diaries from household survey (1-day, home-based social activity)	5 213
S55	Individual single-trip questionnaire (education commute)	597
S56	Trip diaries from household survey (1-day, morning-peak home-work trips only)	6392
S57	Trip diaries from household survey (1–2 day, outbound shopping trips only)	9450
S58	Trip diaries from household survey (home-based trips, sampled to over-represent transit)	4764
S59	Trip diaries from household survey (mode choice analysed at household mobility level)	101053
S60	Individual single-trip questionnaire (mixed purpose)	94
S61	Trip diaries from household survey (1-day)	51910
S62	Individual single-trip questionnaire (education commute)	1484
S63	Trip diaries from household survey (6-weeks, 52 265 trips, 361 individuals, 162 households)	52265
S64	Trip diaries from household survey (1-day, morning-peak trips only)	155016
S65	Trip diaries from household survey (1-day, 7 276 trips, 2 991 individuals, 1 435 households)	7276

Table 9 (*continued*)

No.	Type	N
S66	Unclear household survey	4976
S67	Trip diaries from household survey (grouped into tours: 39 167 home-based tours, 24 396 individuals)	39167
S68	Stated preference - individual panel survey (sampled from 2073 individuals, 7 trips per individual)	8418
S69	Trip diaries from household survey (67 299 trips, unkown individuals, 3 000–3 500 households)	67299
S70	O-D pairs from vehicle tracking data (15000 taxi trips and 15000 bike-sharing scheme trips)	30000

survey data:

- CMAP Travel Tracker Survey, 2007–2008 (Chicago Metropolitan Agency for Planning 2018[b]) - 3 studies [S16; S56; S57]
- CATS Household Travel Survey, 1990 (Chicago Metropolitan Agency for Planning 2018[a]) - [S19]
- Sydney Household Travel Survey (Transport for NSW, 2018) - 2 studies [S32; S69]
- San Francisco Bay Area Travel Survey, 2000 (Metropolitan Transportation Commission 2018[b]) - [S6]
- San Francisco Bay Area Travel Survey, 1990 (Metropolitan Transportation Commission 2018[a]) - [S15]
- Delaware Valley Household Travel Survey, 2012 (Delaware Valley Regional Planning Commission, 2018) - [S60]
- National Household Travel Survey, 2009 (Federal Highway Administration, 2018) - [S67]
- American Travel Survey, 1995 (Bureau of Transportation Statistics, 2018) - [S11]
- Victorian Integrated Survey of Travel and Activity, 2007–2008 (Transport for Victoria, 2018) - [S34]

Three studies make use of academic datasets made public by the authors: S7.1 makes use of the CLOGIT dataset, available with the Ecdat R library (Croissant, 2016; Greene, 2011); S53 uses the SwissMetro dataset (Bierlaire et al., 2001; Bierlaire, 2018); and S63 uses the Mobidrive dataset (Axhausen et al., 2002). Four studies [S29; S40; S41; S52] make use of the partially open LISER PSELL data, which is available on registration (Luxembourg Institute of Socio-Economic Research, 2018). Finally, one study [S70] makes use of openly available bike-sharing and taxi data from the city of Chicago (Divvy Bikes, 2020; City of Chicago, 2020).

Whilst 21 studies make use of open or partially open data, only one study [S48] is identified as making the fully processed data openly available, in the format used for modelling within the paper.

Q2d: Modes in choice-set

Fig. 3 shows a frequency plot of the number of modes considered in each study, which ranges from two to nine. The most common number of modes considered is four, which is used in 18/73 studies.

Five studies have a different number of classes modelled in the classification problem from the number of modes considered. Three studies perform only one-vs-one or one-vs-rest modelling. S11 and S13 both consider three modes, but in both studies the modelling is performed one-vs-rest across the three modes, so that each model considers two different classes. Unlike other studies which use one-vs-rest modelling, the individual models are not combined to create a multiclass classifier in either study. Similarly, S44 considers four modes, but the modelling is performed one-vs-one. As with S11 and S13, the individual models are not combined to create a single multiclass classifier.

Two models jointly model other variables alongside mode choice. S61 jointly considers four modes across two different time-periods (peak/off-peak), therefore modelling a total of eight classes. Similarly, S34 jointly models three modes, three trip purposes, three departure periods, and four distance categories, for a total of 108 classes, 102 of which are observed in the data.

A total of 14 studies use only binary classification. This includes the 11 studies which model only two modes [S12; S17; S18; S37; S42; S45; S51; S54; S55; S62; S70] and the three studies which use one-vs-rest/one-vs-one modelling without combining individual models to a single classifier [S11; S13; S44].

Fig. 4 shows the frequency of each mode/grouping of modes considered in each study. The *car* mode is the most commonly modelled, appearing in 49 studies, followed by *walk* (35 studies) and *public transport* (29 studies). Certain modes either appear individually or grouped. For example, cycling is treated as an independent mode in 25 studies and grouped with walking in nine studies. The grouping of public transport modes cannot be immediately understood from Fig. 4, due to different combinations of groupings being possible. For example, for many studies, rail services are not a viable mode of transport, and so *bus* is the only mode considered. Twenty-nine studies consider all public transport modes under one combined *public transport* mode. Of the 30 studies which consider the independent *bus* mode, 15 include *bus* as the only public transport mode. A total of 19 studies consider two or more separate public transport modes.

Q2e: Modelling of mode-alternatives

In order to understand the impact that the transport network has on mode choice, it is necessary for the dataset to include attributes of the mode-alternatives, e.g. the expected duration and cost of travelling by each mode in the choice-set. These are commonly referred to as Level of Service (LOS) attributes in the literature. For revealed preference data, typically only details of the choice made by the

passenger are recorded. As such, details of the mode-alternatives need to be synthesised and added to the dataset to be included in the modelling.

Of the 68 studies which use revealed preference data, 33 include no attributes of the mode-alternatives in the choice-set [S5; S6; S14; S21; S23; S24; S26; S27; S31; S32; S33; S34; S35; S37; S38; S39; S40; S43; S45; S47; S48; S49; S50; S51; S54; S55; S58; S62; S63; S64; S65; S67; S69]. A further four studies do not list the input features used in the model with enough clarity to deduce whether any attributes of the mode-alternatives are included [S22; S46; S66; S70].

Table 10 shows the relevant features used in the 31 studies which include attributes of the mode-alternatives. The definition of each term is given below:

- *Duration* - journey time from start-point to end-point (including access, transfers etc.)
- *Cost* - Out of pocket cost (e.g. transport fares, Vehicle Operating Costs)
- *Generalised costs* - Combined duration and cost as a single value of disutility, expressed in the unit of currency
- *Vehicle Operating Cost (VOC)* - the mileage dependent costs of operating a vehicle (e.g. fuel, tires, maintenance, repairs, depreciation)
- *In-Vehicle Travel Time (IVTT)* - the duration spent in vehicle/on-board public transport services
- *Out-of-Vehicle Travel Time (OVTT)* - the combined access, egress, transfer, and waiting durations for Public Transport (PT)
- *Access* - The walking duration/distance between the start-point and first public transport access stop
- *Egress* - The walking duration/distance between the last public transport stop and the end-point

Fifteen of the 31 studies which model mode-alternatives do not state the methods used to calculate these attributes [S2; S4; S7.2; S7.3; S10; S12; S13; S15; S17; S20; S25; S28; S30; S57; S59]. Fourteen studies use zonal, time-independent (static) transport models to calculate durations and/or costs [S1; S7.1; S8.1; S8.2; S11; S16; S19; S29; S36; S41; S42; S44; S52; S61]. One study [S8.2] additionally makes use of a time-dependent public transport model to calculate transfer and combined access/egress durations for the PT route at the time of departure. Finally, two studies [S56; S60] make use of an online directions service to generate trip durations.

Q2f: Input features dependent on output choice

In order to be used as a valid predictive model, model input features must be independent of the output choice. Features which are dependent on the choice, e.g. the recorded trip duration (which is dependent on the mode taken) cannot be known until the trip is made, and so cannot be used for prediction.

A substantial proportion of studies (19/73) include input features which are related to the output choice, either directly or indirectly.

Ten studies [S6; S23; S31; S40; S44; S45; S54; S58; S62; S63] include the recorded travel duration of the selected mode as an input feature. Five of these studies also include the trip distance [S31; S40; S44; S58; S63], which would allow the classifier to infer the speed of the mode-selected. A further two studies [S6; S31] additionally include the reported cost of the selected mode.

Two studies [S16; S32] implicitly include the reported duration by including both the reported departure time and arrival time in the feature vector.

Three studies [S7.1; S7.2; S7.3] implicitly include the selected mode in the input feature vector by labelling attributes of the *selected* mode and best *alternative* mode. For example, one node in the DT for S7.2 separates trips between those made by *Auto* and those made


Fig. 3. Frequency bar chart of number of modes considered in each study in review.

by *Metro* on the basis of whether the cost of the *selected mode* is greater than or equal to 1.6 euro.

Two studies use different definitions of duration in the mode-alternative attributes for the selected mode. S52 uses the reported duration from the survey as the driving duration if the trip is made by car and uses the driving time predicted by a static zonal transport model if the trip is not made by car. S60 similarly uses the reported duration for the selected mode, and the duration as predicted by the Google Directions Application Programming Interface (API) for all other modes. In both cases, treating the selected mode differently to the unchosen alternatives may cause leakage of the selected mode into the input feature vector.

Finally, two studies [S37; S43] include survey questions on reasons for not taking a particular mode in the input feature vector.

As with the modelling of mode-alternatives, four studies [S22; S46; S66; S70] provide insufficient detail of the modelling process to determine whether input features are included which are dependent on the output choice.

Q2g: Hierarchical data

As shown by Q2a, 36 studies make use of trip diary data. Household trip survey data has an inherent hierarchical structure: households are made up of multiple people, each of whom make multiple tours, in which there are multiple legs or trips. Elements within the same groups in the hierarchical structure may show interdependency. This hierarchical structure arises from the specific nature of how trip diary data is collected, and introduces strong correlations which can be observed in the data. Formally, three levels of hierarchy can be considered (each with examples of how the structure could cause interdependency):

- *Household-Person (H–P)* - e.g. multiple members of a household travelling together therefore all travelling by the same mode, one person using the only vehicle in a household meaning that others cannot use that vehicle, all members of a household sharing a tendency to/not to travel by a particular mode, etc.
- *Person-Tour (P-T)* - e.g. individual showing a tendency to/not to travel by a particular mode, individual not being able drive/cycle for all tours due to a vehicle/bike not being available to them on the survey date, individual having a season ticket and therefore being more likely to travel by public transport, etc.
- *Tour-Trip (T-T)* - e.g. return trip being highly likely to be made by the same mode as the outbound trip, vehicle/bike not being available for onwards travel as it was not used for first leg (trip) in tour, vehicle/bike needing to be used for onwards travel as it was used for first leg (trip) in tour and cannot be left behind, etc.

Individual survey trip diaries do not have a household-person grouping, leaving person-tour and tour-trip groupings.

Whilst they do not use trip diaries, household surveys where multiple members of the same household are interviewed to extract commute (as in S29 and S41) also contain a household-person hierarchical structure.

As well as RP surveys, panel SP data, where each individual provides multiple choices in a survey, also has a Person-Tour hierarchy.

The details of all studies which use data which has or may have a hierarchical structure are shown in Table 1. This includes the studies with datasets of unknown nature, which may be hierarchical [S1; S10; S12; S17; S22; S31; S38; S46; S66].

Many of the studies which make use of trip diary data sample the data in a way which removes all/part of the hierarchical structure, e.g. by sampling only outbound trips (removes tour-trip hierarchy), or by sampling only trips made by one member of a household

Fig. 4. Frequency bar chart of individual modes/grouping of modes in each study in review. The 'Other' category groups all modes/combinations of modes with less than three occurrences across all studies.

(removes household-person hierarchy). This sampling is also presented in Table 11.

Table 11 shows the levels present in the input dataset (after any sampling/processing) for all studies which make use of hierarchical data. Whilst S16 uses a tour-based analysis, it still predicts mode choice for individual trips, and so the Tour-Trip hierarchy in the data is still present. In total, there are 45 studies which use hierarchical data, or data which may be hierarchical, after sampling/processing. This includes 13 studies which use complete, unsampled trip diaries [S16; S30; S32; S35; S40; S42; S44; S47; S48; S60; S61; S63; S65; S69]. Note that using hierarchical data is not an issue in itself, as long as appropriate sampling is used for validation. The sampling used in these studies (and its associated implications) is therefore discussed in Q3b.

4.3.1. Model datasets - limitations

Five limitations are identified in relation to the datasets used to investigate mode choice. Two limitations are technical: (i) studies not including any attributes of the mode-alternatives, (ii) studies using input features dependent on output choice; and three limitations are general: (i) not describing the dataset and modelling process in sufficient detail, (ii) the lack of relevant, openly available datasets including mode-alternative attributes, (iii) not considering sampling of the data from the population.

Note that using hierarchical data is not an issue in itself, as long as appropriate sampling is used for validation. This is therefore discussed in Section 4.3.

Two technical limitations are identified related to datasets. Q2f identifies 33 studies which include no LOS attributes of the mode-alternatives in the choice-set, and a further four studies which do not list the input features used in the model with enough detail to be able to determine whether any attributes of the mode-alternatives are included. In order to model the impact that the transport network has on mode choice, it is necessary for the feature vector to contain attributes of the mode-alternatives, e.g. the expected duration and cost of travelling by each mode in the choice-set. As significant correlations between attributes of each mode-alternative and mode choice are likely to exist, not including these variables in the feature vector will result in models with lower predictive performance. Additionally, for statistical RUM models, omitting relevant predictors (features) in the input results in endogenous errors in the parameters of the remaining variables (Train, 2009, Chapter 13). This can cause biased, inconsistent estimates of these parameters, which may be important for explaining behaviour (e.g. VoT). Finally, when using the choice model for simulation of future trips under unknown conditions (e.g. in an Agent Based Model (ABM)), the impacts of changes to the transport network on mode choice cannot be modelled if attributes of the mode-alternatives are not included in the feature vector. These studies therefore do not allow for modelling the impacts of changes to the transport network on the mode choice decisions made by an individual.

Of the studies which do model mode-alternatives, the majority generate LOS variables from static zonal graphs. This means that they do not capture the highly granular spatial and temporal variability of conditions on a transport network.

Q2f identifies 19 studies which include input features which are related to the output choice. These features cannot be known in advance of a trip being made, and so this prevents these models from being used in a predictive context. Additionally, input features which are directly and explicitly dependent on class membership, e.g. travel speed being dependent on travel mode, may be highly correlated with the class membership. As such, this will result in better *apparent performance* of the model than could be achieved using only valid independent variables (i.e. the performance of the model will be overestimated through *data leakage*). As with omitting the mode-alternative LOS variables, including input variables which are dependent on the output in a statistical model (RUM) can

Table 10

Attributes of mode-alternatives in selected studies in review. Unless stated otherwise, each attribute is a duration. *PT*=Public Transport, *IVTT*=In-Vehicle Travel Time, *OVTT*=Out-of-Vehicle Travel Time, *VOC*=Vehicle Operating Cost.

No.	Duration	Cost	Other	No.	Duration	Cost	Other
S1			IVTT and route distance (Each mode)	**S20**	✓	✓	
S2	✓	✓		**S25**			Generalised costs (Each mode)
S4	✓	✓	Access (Bus)	**S28**	✓	✓	
S7.1	✓	✓	IVTT (PT)	**S29**			Generalised costs (Each mode)
S7.2	✓	✓		**S30**		✓	IVTT and OVTT (Each mode)
S7.3	✓	✓	IVTT (PT)	**S36**	✓		Access distance (PT)
S8.1	✓	✓	Transfer, access/egress (PT)	**S41**			Generalised costs (Each mode)
S8.2	✓	✓		**S42**			IVTT, transfer, speed, directness (PT)
S10			OVTT (Bus)	**S44**	✓		
S11	✓	✓		**S52**	✓	✓	
S12	✓	✓		**S54**	✓	✓	
S13			Duration ratios (each mode)	**S57**	✓		Access & egress distance (PT)
S15	✓	✓		**S58**	✓	✓	
S16	✓	✓		**S60**		✓	IVTT and OVTT (Each mode)
S17		✓	Access, egress, IVTT (PT)	**S61**	✓		
S19			Duration, VOC (Drive), IVTT (train)				

introduce endogeneity through reverse causality (Train, 2009, Chapter 13). This can also cause biased, inconsistent estimates of model parameters. Again, a further two studies provide insufficient detail of the modelling process to be able to determine whether any input features which are dependent on the output choice are included.

Of the two technical limitations related to datasets, using input features dependent on output choice a *pitfall* that is likely result in incorrect conclusions being drawn from the modelling results. Conversely, not including any attributes of the mode-alternatives is an *area for improvement*, as doing so is likely to improve the performance of the model.

The discussion of the research question also highlights four general limitations. Firstly, multiple studies do not describe the dataset and modelling process in sufficient detail for the required information for the systematic review to be extracted. This is problematic for repeatability of the mode choice experiments implemented in these studies, particularly when there is such large variation in the methodologies used in each study. In order to ensure repeatability of the results, methodologies should be recorded in detail, and where possible, data and code should be made available.

The discussion also highlights the need for relevant, openly available datasets including mode-alternative attributes. There exist several openly available, large datasets for investigating passenger mode choice. Of the 16 studies which use datasets with greater than 10 000 entries, 11 make use of openly available datasets [S11; S14; S16; S32; S34; S40; S48; S63; S67; S69; S70]. However, only two of these studies [S11; S16] add mode-alternative information to these datasets, and the processed dataset is not openly available for either study. As mentioned, only the processed dataset for S48 is openly available, and this dataset does not include any mode-alternative attributes. For an example of a large, openly available dataset with mode-alternative attributes see Hillel et al. (2018).

Finally, no studies checked the representivity of the sample in the dataset with respect to the target population, or discussed how to correct for sampling biases in forecasting. When using a model for forecasting, it is essential to consider the bias in the sample, for example for accurate predictions of market shares.

Table 11
Details of hierarchies in datasets in relevant studies in review, after sampling/processing. *H–P*=Household-Person, *P-T*=Person-Tour, *T-T* = Tour-Trip.

No.	H–P	P-T	T-T	Sampling	No.	H–P	P-T	T-T	Sampling
S1	?			Work trips only, sampled from larger survey	**S39**	?	?	?	Education trips only
S3		✓		None (SP - Multiple choices per person)	**S40**		✓	✓	None (complete trip diary, individual)
S5		✓		None (complete activity diary, individual)	**S41**	✓			None (Household survey)
S6	?	?		Outbound work trips only (2-day)	**S42**		✓	✓	None (complete trip diary, individual)
S9		✓		None (SP - Multiple choices per person)	**S43**				Mobility of head of household only
S10	?	?	?	Unclear data	**S44**	✓	✓	✓	None (complete trip diary, household)
S11				Mobility of similar clusters	**S46**	?	?	?	Unclear data
S12	?	?	?	Unclear data	**S47**		✓	✓	None (complete trip diary, individual)
S13				Commute patterns from individual survey	**S48**	✓	✓	✓	None (complete trip diary, household)
S15	?			Outbound work trips only	**S51**				Household mobility only
S16	✓	✓	✓	None (complete trip diary, household)	**S52**				Commute patterns from individual survey
S17	?	?	?	Unclear data	**S53**		✓		None (SP - Multiple choices per person)
S19	✓			Morning home-work trips only	**S54**	✓	?		Home-based social trips only
S21	?			Outbound work trips only	**S56**	✓			Morning home-work trips only
S22	?	?	?	Unclear data	**S57**	✓	?		Outbound shopping trips only
S24	?	?	?	Random sampling from larger survey	**S58**	✓	✓		Home-based trips only
S26	?	?	?	Random sampling from trip diaries	**S59**				Household mobility only
S27	?	?	?	Random sampling from larger survey	**S61**	✓	✓	✓	None (complete trip diary, household)
S29	✓			None (Household survey)	**S63**	✓	✓	✓	None (complete trip diary, household)
S30	✓	✓	✓	None (complete trip diary, household)	**S64**	✓	✓		Morning peak trips only
S31	?	?	?	Unclear data	**S65**	✓	✓	✓	None (complete trip diary, household)
S32	✓	✓	✓	None (complete trip diary, household)	**S66**	?	?	?	Unclear data
S33	✓	✓		Outbound trips only	**S67**	✓	✓		Tours from household trip diary
S34	✓			First trip in day only	**S68**		✓		None (SP - Multiple choices per person)
S35	✓	✓	✓	None (complete trip diary, household)	**S69**	✓	✓	✓	None (complete trip diary, household)
S38	?	?	?	Unclear data					

4.4. How is model performance determined?

The following sections discuss the techniques used to determine model performance in the 73 studies in the review, focusing in turn on the validation method, the sampling method, and the performance metrics used.

Q3a: Validation method

The validation method most commonly used in the studies is holdout validation (non-repeated), which is used in 50 studies. Train-test splits range from 23:77 to 91:9, but the most commonly used splits are 70:30 (12 studies), and 80:20 (12 studies).

Seven studies use repeated holdout validation: S26 runs 50 repetitions of a 70:30 split, S37 runs 10 repetitions of a 75:25 split, S53 runs 10 repetitions of a 70:30 split, S54 runs 3 repetitions of a 75:25 split, S60 runs 100 repetitions of a 75:25 split, and finally S29 and S41 run 100 repetitions of a 60:40 split. Confusingly, S26 only shows the results for both the train and validation data combined, averaged over the 50 runs.

k-fold cross-validation is used in seven studies. Four studies use 10-fold cross-validation [S34; S48; S57; S61], two studies use 5-fold cross-validation [S52; S67], and one study [S70] uses 3-fold cross-validation. As well as 10-fold cross-validation, S57 also performs holdout validation (60:40 split).

Four studies use different validation techniques for different models. Whilst they all use in-sample validation for the logit models, S9 uses 80:20 holdout validation for the neuro-fuzzy multinomial logit model; S7.2 uses 60:40 holdout validation for the ANN, DTY, and discriminant analysis models; and S58 uses Out-Of-Bootstrap (OOB) error for the RF model. Furthermore, S70 uses 3-fold cross validation for the majority of the ML classification techniques but uses holdout validation for the ANN models.

Five studies use in-sample validation for all models [S5; S14; S23; S45; S59].

Three studies do not state the validation method used [S16; S42; S50].

Finally, one study [S55] does not perform any validation of the final model, instead using 5-fold cross-validation for hyperparameter selection (see Q4b), and then extracting behavioural indicators from the structure of the final model.

Q3b: Sampling method

Of the 45 studies which use hierarchical data, or data which may be hierarchical, none mention the use of grouped (by household or individual) sampling. This includes all 13 studies which make use of complete, unsampled trip diaries. Furthermore, two studies which make use of trip diaries [S16; S42] do not state which validation technique is used at all (see Q3a).

All studies which perform out-of-sample validation appear to use random sampling (either stated explicitly or assumed).

Only two studies [S3; S64] test models on data collected separately from, or after, the training data (*external validation*). In S3, each city-specific model is additionally validated on the data from the other city (i.e. the model trained on Melbourne is validated on the data from Sydney and vice-versa). In S64 the model estimated on travel survey data collected in 2008 is tested on the data collected in 2013.

Q3c: Performance metrics

The performance metrics used for model validation in each study are shown in Table 12. Note, that Table 12 and the discussion in this section only considers the metrics used in the studies in the review. There are many other relevant metrics which can be used to evaluate classifier performance, though if they are not used in the studies they are not discussed here.

The first four columns of Table 12 show discrete metrics, where each trip is assigned to the mode with highest predicted probability. This is used to produce the *confusion matrix*, from which the other metrics (accuracy, recall, and the mode-shares) are derived. Three further discrete metrics (which can also be calculated from the confusion matrix) are not shown in Table 12, as they are only used in one or two studies each. Precision is used in S16 and S70, specificity is used in S13, and F1-score is used in S70.

Metrics which evaluate probabilistic classification (i.e. which evaluate the probability distributions generated by classifiers, and not the discretised maximum probability classes) are grouped together in Table 12. Seven different probabilistic metrics are used in the 10 different studies: percent clearly right ($p_i > 0.9$)/clearly wrong ($p_i < 0.1$)/unclear ($0.1 <= p_i <= 0.9$) (i.e. where different probability thresholds are used to classify the confidence of the prediction) [S4; S8.1; S8.2], Arithmetic Mean Probability of Correct Assignment (AMPCA) (referred to as fitting factor in S4, S8.1, and S8.2; and average probability of correct assessment in S41), Mean Squared Error (MSE) [S4; S8.1; S8.2; S46], simulated mode shares [S4; S8.1; S8.2; S53], Receiver Operating Characteristic (ROC) curves [S38; S49; S61], log-likelihood [S52; S53], Bayesian Information Criterion (BIC) [S52], and Expected Simulation Recall (ESR) [S41].

Table 12 does not show the metrics used in three studies. As explained in Q3a, S55 does not perform validation of the final model, instead extracting behavioural indicators (sensitivities) from the structure of the final model. Note that three further papers [S2; S9; S68] also extract behavioural indicators from the model structure (see Q5b), however these papers validate the models using the confusion matrix/accuracy, which are included in Table 12. S11 performs regression on the total number of trips performed by each mode within a cluster, and so uses regression-based metrics (MSE and average relative variance of regression). S3 uses three metrics: *predicted share less observed share*, *weighted percent correct*, and *weighted success index*. However, no definitions for the performance metrics are provided in the paper, and so it cannot be determined if the metrics are discrete or probabilistic.

In total, 60 of the remaining 70 studies use only discrete classification metrics (including three which extract behavioural indicators

from the probability distributions) and 10 studies use a combination of probabilistic and discrete classification metrics. Of the studies which use only discrete classification metrics, 35 make use of LR models.

Fourteen studies use only one performance metric: accuracy is used as the sole metric in 9 studies [S10; S17; S25; S34; S35; S42; S47; S50; S54; S58; S62], recall per mode in S45, and the confusion matrix in S2.

4.4.1. Model performance estimation - limitations

Four technical limitations are identified in relation the model performance estimation techniques used in the studies: (i) studies using inappropriate validation schemes, (ii) studies using incorrect sampling methods for hierarchical data, (iii) studies not performing

Table 12

Summary of performance metrics used for validation in each study in review. **Acc** = Accuracy, **Rec** = Recall, **CM**=Confusion Matrix, **MS** = Mode Shares (Discrete), **Pro** = Probabilistic metric.

No.	Acc	Rec	CM	MS	Pro	No.	Acc	Rec	CM	MS	Pro
S1	✓	✓				S35	✓				
S2			✓			S36	✓	✓			
S3	–	–	–	–	–	S37	✓	✓			
S4	✓	✓			✓	S38	✓	✓	✓		✓
S5	✓	✓		✓		S39	✓		✓	✓	
S6	✓	✓	✓	✓		S40	✓	✓			
S7.1		✓	✓			S41			✓		✓
S7.2	✓	✓	✓			S42	✓				
S7.3		✓	✓			S43	✓	✓			
S8.1	✓	✓			✓	S44	✓		✓		
S8.2	✓	✓			✓	S45		✓			
S9	✓		✓			S46	✓				✓
S10	✓					S47	✓				
S11	–	–	–	–	–	S48	✓	✓	✓		
S12	✓		✓			S49	✓		✓		✓
S13	✓	✓		✓		S50	✓				
S14	✓		✓			S51	✓	✓	✓		
S15	✓	✓	✓			S52	✓		✓		✓
S16	✓	✓				S53	✓			✓	✓
S17	✓					S54	✓	✓	✓		✓
S18	✓		✓			S55	✓				
S19	✓	✓	✓			S56	–	–	–	–	–
S20		✓		✓		S57	✓	✓			
S21	✓		✓			S58	✓		✓		
S22	✓	✓	✓			S59	✓				
S23	✓	✓				S60	✓		✓		
S24	✓	✓		✓		S61	✓	✓		✓	
S25	✓					S62	✓				
S26	✓	✓	✓			S63	✓	✓			
S27	✓	✓	✓			S64	✓	✓	✓		
S28	✓		✓	✓		S65	✓			✓	
S29	✓	✓	✓			S66	✓	✓	✓		
S30	✓	✓		✓		S67		✓	✓		
S31	✓	✓	✓			S68	✓				
S32	✓			✓		S69	✓		✓		
S33	✓			✓		S70	✓	✓			
S34	✓					**Sum**	63	38	33	13	10

external validation, (iv) studies using only discrete metrics.

Q3a identifies 12 studies which make use of inappropriate validation schemes. This includes five studies which use in-sample validation [S5; S14; S23; S45; S59], four studies which use different validation techniques for different models being tested [S7.2; S9; S58; S70], and three which do not state the validation method being used [S16; S42; S50].

In-sample validation uses the same data to fit and validate the model and can be interpreted as using the *train-error* to estimate model performance. As such, it presents only the explanatory power of the model, i.e. the ability of the model to replicate the training data, and not the predictive performance. This is discussed in detail by Shmueli (2010). If a model has high *variance*, it can overfit to noise in the data during model fitting, without generalising to valid correlations between the input and output. This will result in in-sample validation overestimating the predictive performance. Without testing the model on out-of-sample data, there is no way to assess whether overfitting has occurred. Additionally, due to the nature of the *bias-variance tradeoff* (Hastie et al., 2008, Chapter 2), a classifier will tend to fit partially to noise in the data, even if it does not overfit. As such, the train-error will tend to overestimate predictive performance, even for well specified models which do not overfit.

Formal validation of a model on data separate training data is essential to ensure ML models have generalised to the training data without overfitting. As such, in-sample validation is an inappropriate validation scheme. Furthermore, in order to make valid comparisons between performance estimates of different models, the same validation method must be used for all models. Otherwise, any apparent differences in performance may be due to differences in the respective validation schemes.

Q2g identifies 29 studies which make use of hierarchical data, and a further 16 which makes use of data which may be hierarchical. As identified by Q3b, none of these studies sample validation sets or folds grouped by individual or household. As such, trips from the same group (household/person/tour) will occur in both test and training data. These trips inherit correlated features from these groupings, which can allow for data-leakage and overfitting.

There are valid hierarchies in datasets which can be relevant to the modelling scenario. For example, a modeller would be interested if students (socio-economic group) show a tendency towards cycling (correlation), or if trips made at the weekend (temporal grouping) were less likely to be made by public transport (correlation). In both these cases, the hierarchies (groups) are general, and described by information in the feature vector. As such, these correlations are likely to be constant across the training data and future unknown trips, and so are relevant to the modelling scenario.

Conversely, the hierarchies identified by Q2g are not representative of the population (and instead are a feature of our data sampling). As such, these hierarchies are not relevant to the modelling scenario, and will boost the *apparent performance* of the model, whilst in reality causing it to perform worse on true unseen data.

This is particularly problematic for the 13 studies which use complete trip diary data [S16; S30; S32; S35; S40; S42; S44; S47; S48; S60; S61; S63; S65; S69]. Many of these trip diaries are multi-day, compounding the issue. Notably, S48 uses a six-day trip diary (average 3.3 trips per person), S47 uses sets of GPS trips logged over four months (average 58.4 trips per person), and S63 uses a six-week travel diary (average 144.8 trips per person). This problem is not unique to mode choice modelling applications. Saeb et al. (2017) conduct a review of sampling methods in studies using ML to make clinical predictions from smartphone or wearable technology data. They review studies which use hierarchical data, where there are multiple *records* for each individual *subject*. They find that of the 62 of the studies included in the meta-analysis, 28 (45%) use inappropriate *record-wise* sampling, instead of *subject-wise* sampling.

Q3b identifies that only two of the studies reviewed use *external validation*, where the model is validated on data collected separately from, or after, the training data [S3; S64]. External validation using future data is the only possible method of directly simulating the use case for a mode choice model, of predicting future, unknown trips. External validation can also identify issues with data-leakage, overfitting, and incorrect validation schemes, e.g. the incorrect sampling methods for hierarchical data, as highlighted by Q2g.

Finally, Q3c identifies that the vast majority of studies (60/73) use only discrete metrics to assess model performance, where each trip is assigned to the mode with the highest predicted probability. This includes 35 studies which assess LR using only discrete classification metrics, despite LR being a statistical technique intended to generate probability distributions. In total, only six studies make use of *strictly proper* continuous scoring metrics, log-likelihood [S52; S53] and MSE [S4; S8.1; S8.2; S46] (Gneiting and Raftery, 2007). Note that the log-likelihood (also known as logarithmic score and Cross-Entropy Loss (CEL)) can be normalised by dividing by the sample size. This can allow for comparison between samples of different sizes.

There are a number of issues with using only discrete metrics to assess choice prediction. Firstly, discretising the classification by assigning each observation to the highest probability class is not a proper use of output choice probabilities, and will likely result in non-representative mode-shares in imbalanced data. Mode choice data is inherently imbalanced, i.e. there are likely more trips made by some modes (e.g. car, walking) than others (e.g. cycling). By assigning each prediction to the class with highest probability, the less frequent classes will be under-represented in the predicted outcomes, and the more frequent classes will be over-represented. For example, consider a biased random coin flip, where heads is 60% likely to occur, and tails occurs with 40% probability. The best possible predictive classification model will predict these outcomes at their respective probabilities for each coin flip event. However, assigning the highest probability class for the prediction will result in heads always being predicted (and never tails) as heads is always more likely than tails. This clearly results in non-representative class shares. Non-representative mode-shares are unacceptable for mode choice models, where the mode-shares are a crucial model output. Where discrete predictions are needed from a probabilistic model, the assignment should be simulated by drawing the predicted labels from the output probability distributions. This results in representative mode shares.

Similarly, by generating a discrete class for each observation, mode choice is treated as a deterministic instead of a stochastic process. As such, it is assumed that mode choice is constant under the same set of conditions and socio-economic characteristics. In

reality, the model does not contain all information describing the choice, and passengers have a degree of intra-heterogeneity. As such, a passenger's choice can be considered as being drawn randomly from a probability distribution given the observed features in the model (as with the coin-flip example). We define this distribution as the *true model*, which we aim to replicate with the classification model. In order to account for this stochastic heterogeneity in simulation, the predicted mode choice should be drawn from a probability distribution. The metric used to assess model performance should therefore represent how well the predicted probability distributions fit the data.

Additionally, discrete metrics do not assess how right or wrong model predictions are. For example, when using discrete metrics, the contribution to the model's score for a trip where a binary classifier predicts the selected mode at 1% probability is the same as that for a trip where the classifier predicts the selected mode at 49% probability. Analysing the probability distributions presents information on where the model performs well or poorly.

Finally, by taking the maximum of the class probabilities, discrete predictions and the associated metrics are discontinuous. This results in discrete metrics having an expected score which is not differentiable or strictly convex. Additionally, accuracy and other discrete metrics are not *strictly proper* scoring rules, and as such do not have unique maximums (Gneiting and Raftery, 2007). This makes discrete metrics poor metrics to use during model fitting, particularly where a continuous gradient is required (e.g. gradient descent).

Note that the use of discrete metrics alongside strictly proper scoring rules as an easily interpretable indication of performance to be compared between studies is not considered as a limitation in this study. Instead, it is only considered as a limitation if a paper uses *only* discrete metrics, with no probabilistic metrics.

Of the four technical limitations related to model performance estimation, three represent pitfalls (using inappropriate validation schemes, using incorrect sampling for hierarchical data, and using only discrete metrics), and one represents an area for improvement (not performing external validation).

4.5. How are optimal model hyper-parameters selected?

Hyper-parameters are parameters of classification algorithms which are used to regularise the model during model training. The selected hyper-parameters impact the bias and variance of the fitted model. This section discusses the techniques used to optimise model specifications and hyper-parameters for conventional ML classification algorithms (ANNs, DTs, EL, SVMs). The 14 studies which do not use at least one these algorithms are therefore omitted from this section of the review [S9; S10; S12; S14; S18; S19; S25; S28; S30; S33; S39; S40; S42; S52].

The following sections review the remaining 49 studies, focusing in turn on the hyper-parameter search method, the hyper-parameter validation method, and the hyper-parameter validation data.

Q4a: Hyper-parameter search method

Of the 59 studies which use at least one conventional ML algorithm, 11 do not mention hyper-parameter values at all within the paper [S7.1; S7.2; S7.3; S23; S27; S32; S35; S45; S46; S50; S58]. A further 10 studies either state hyper-parameter values without explanation [S2; S13; S34; S36; S43; S61; S67; S69], or state that they use default values [S24; S51].

This leaves 38 studies which use some form of hyper-parameter optimisation. Fifteen studies [S1; S3; S4; S5; S11; S16; S21; S37; S54; S56; S57; S60; S62; S64; S68] perform a manual search, or trial and error, in order to identify model parameters. Of these, S16 searches only for the kernel function in an SVM and uses default values for all other parameters and models, and S1 searches for the number of neurons in a single test layer, again using default values for other parameters.

Eleven studies [S15; S17; S20; S22; S26; S29; S31; S41; S44; S49; S59; S66] specify an MLP with a single hidden layer and perform a linear search on the number of neurons in that layer. With the exception of S15, which performs a grid search for the SVM parameters (γ and C), default values are used for all other parameters of all models.

One study [S44] uses a repeated linear search, firstly on the loss-weight ratio of the two classes in each model, and secondly on the number of features used.

Seven studies [S6; S8.1; S8.2; S48; S53; S63; S65; S70] make use of a grid search to find optimal hyper-parameters. This includes S65, which optimises the RF model using a grid-search, but uses default values for all other parameters.

One study [S38], tests two different search strategies in order to find optimal SVM parameters (γ and C): grid search and genetic algorithms. The study finds that whilst the two methods find optimal solutions with similar accuracies, the genetic algorithm finds the solution with the lower penalty parameter (C), and so is preferred.

One study [S55] uses the *early stopping* method of GBDTs, where DTs are sequentially added to the ensemble until the out-of-sample predictive performance stops improving.

Finally, one study [S47] states that cross-validation is used to select model parameters but does not state the search method used.

Q4b: Hyper-parameter validation method

Of the 38 studies which use some form of hyper-parameter optimisation, 13 do not state the validation method used to determine optimal values [S1; S11; S16; S20; S22; S26; S31; S37; S54; S57; S62; S64; S70].

Eleven studies [S3; S4; S6; S15; S29; S49; S53; S56; S59; S66; S68] use holdout validation. Nine studies [S5; S21; S38; S44; S47; S48; S57; S60; S63] use *k*-fold cross-validation. One study [S41] uses repeated holdout validation. One study [S17] uses in-sample

validation. One study [S65] uses the OOB error of the DTs in the RF.

Finally, two studies [S8.1, S8.2] use a complex multi-criteria assessment, involving relative performance on both the calibration and validation data.

Q4c: Hyper-parameter validation data

Of the 38 studies which use some form of hyper-parameter optimisation, 19 do not state the data used for hyper-parameter validation [S1; S3; S11; S16; S20; S22; S26; S29; S31; S37; S38; S47; S54; S57; S62; S63; S64; S66; S70].

Of the nine studies which use *k*-fold cross-validation to test hyper-parameter performance, two use only the training data [S15; S21], one uses a random subset of 43% of the data [S48], three use all of the data [S5; S55; S60], and three do not state the data used (included above). The study which uses repeated validation also uses all of the data [S41].

Of the 11 studies which use holdout validation, three use the data reserved for model testing [S4; S53; S56], two use only the train data, dividing it into a new test and train fold [S6; S21], three uses a separate validation sample which is not used for model testing or training [S49; S59; S68], and three do not state the data used (included above).

Finally, the two studies which use the multi-criteria assessment [S8.1; S8.2] use both the train and test data.

4.5.1. Model optimisation - limitations

Four limitations are identified in relation the model optimisation techniques used in the studies. Three limitations are technical: (i) studies not performing any type of hyper-parameter optimisation, (ii) studies not using rigorous hyper-parameter search schemes, (iii) studies optimising hyper-parameters on validation data; and one is general: not presenting model hyper-parameters used within the study.

Three technical limitations are identified by the attributes related to hyper-parameter optimisation collected in the systematic review. Of the 59 studies which use one or more conventional ML models to investigate mode choice, Q4a identifies 21 studies which do not perform any type of hyper-parameter optimisation. This includes 11 which do not state hyper-parameter values at all, and 10 which use default values or provide values without explanation. Model performance is highly dependent on chosen hyper-parameter values. Additionally, optimal hyper-parameter values are highly task dependent, and will vary for different datasets, metrics, scenarios, etc. Using default hyper-parameter values, or values from previous studies with different modelling scenarios or data, is therefore likely to result in sub-optimal hyper-parameters being used, and the resultant model will perform worse than the optimised model. If the hyper-parameters of each classifier have not been optimised, it is not possible to make valid comparisons between the respective algorithms, as any difference in model performance may be due to better hyper-parameter values selected for one algorithm than another.

Q4a also identifies that no studies use a fully rigorous hyper-parameter search method. Many studies use inconsistent search methods, only searching over one parameter within one model (e.g. number of neurons in a hidden layer), whilst leaving all others with default values. Optimising only the parameters for only a subset of classifiers being compared will tend to improve the performance of those classifiers over those which have not been optimised. Additionally, the search space should cover all dimensions of the hyper-parameter space, otherwise optimal values are unlikely to be found. Whilst certain hyper-parameters may have little/no effect on model performance, there is no way to determine this unless they are tested.

Additionally, search schemes should be used which maximise the probability of finding optimal hyper-parameters in an unbiased manner. Only one study uses an automated sequential search (genetic algorithm in S38) to optimise model hyper-parameters, the rest either using a pre-specified search space (linear search/grid search) or manual search/trial and error.

The primary advantages of manual search are its simplicity and the ability to use the modeller's intuition (from previous trials and similar classification tasks) to influence subsequent guesses. However, manual search presents both high potential for the introduction of bias, and difficulty in reproducing results. Additionally, as the search is manual and cannot be parallelised, it practically limits the modeller to a small number of trials in *S*.

Grid-search predefines a set of candidate values for each hyper-parameter and use them to define a search space *S* containing each unique combination of values. Grid-search can be both automated and parallelised, and therefore enables a greater set of candidate values to be searched than with a manual search. However, grid-search is unable to learn from previous evaluations, and so spends a lot of time evaluating candidate values which are unlikely to perform well. Additionally, the same values for each hyper-parameter are repeated for each dimension of the search, limiting the likelihood of evaluating the optimal value for each hyper-parameter. As such, grid-search is highly inefficient for hyper-parameter selection and has been shown to perform poorly in practice at finding optimal hyper-parameter values compared to other search schemes, including random search (Bergstra and Bengio, 2012).

Finally, Q4c identifies eight studies which include the validation data in the hyper-parameter search [S4; S5; S8.1; S8.2; S41; S53; S56; S60], as well as 19 which do not state the data used [S1; S3; S11; S16; S20; S22; S26; S29; S31; S37; S38; S47; S54; S57; S62; S63; S64; S66; S70]. Fitting hyper-parameters to the holdout validation data allows the model to select optimal hyper-parameters specifically for that data. In other words, this presents the potential for the model to fit to the validation data using the hyper-parameters (*data leakage*). This will upward bias the performance estimate over that which would be achieved with previously unseen data. This is explored by Varma and Simon (2006), who show that cross-validation provides an upward biased estimate of true performance if it is used for model optimisation.

As discussed, validating a model on previously unseen data is an essential step in predictive modelling. Holdout validation data should not be seen by the model at any time during model development (including hyper-parameter optimisation) until the testing of the finished model.

Of the three technical limitations related to model optimisation, one represents a pitfall (optimising hyper-parameters on validation data), and two represent areas for improvement (not performing hyper-parameter optimisation, and studies not using rigorous hyper-parameter search schemes).

The discussion of Q4c also highlights one general limitation, that studies do not report the model hyper-parameters and hyper-parameter selection schemes with sufficient detail. As with the details of methodologies in Q2, this is problematic for repeatability of the mode choice experiments implemented in these studies. Hyper-parameter values and selection schemes should be recorded in detail in order to ensure repeatability of the studies.

4.6. How are the final models analysed?

This section discusses how the finalised models (i.e. after training, optimisation, and validation) are analysed, focusing in turn on statistical testing; and the extraction of behavioural indicators.

Q5a: Statistical testing

Across all 73 studies, only four [S26; S48; S53; S60] conduct any analysis of the uncertainty or distribution of model performance. S48 uses 10-fold cross validation to estimate the accuracy of seven different classifiers. Firstly, the study uses a Kruskal-Wallis test at a 5% significance level to test the null hypothesis that the performance estimates of all classifiers tested are not significantly different from one-another. Secondly, a two-sided Wilcoxon rank-sum test is applied pairwise between the classifiers to test whether different pairs of classifiers are significantly different from each other.

Three studies [S26; S53; S60] estimate the standard deviation of the metrics (accuracy in S26 and S53, and accuracy and recall in S60) across each run of *k*-fold cross-validation/repeated holdout validation. These estimates of standard deviations are not used to form any formal significance tests in these studies.

Q5b: Extraction of behavioural indicators

As ML classifiers do not have an underlying behavioural model, it is not straightforward to extract behavioural indicators (e.g. VoT and choice elasticities). Several papers include details of the ML models' structural information, including DT structure, feature importances, decision rules from Rough Set Models (RSMs), etc. However, this is not equivalent to the behavioural indicators that can be obtained from RUMs.

Overall, only four papers [S2; S9; S55; S68] attempt to extract standard behavioural indicators from the ML classifiers. These four papers all perform elasticity or sensitivity analysis: S2 calculates aggregate point elasticities by modifying the variable of interest for all observations in the dataset and rerunning the model, whilst the remaining three papers conduct disaggregate analysis at the mean values for the other variables (sensitivities in S9 and S57 and probability derivatives in S68). S68 additionally calculates VoT estimates from the probability derivatives for travel time and cost.

Five further papers [S12; S28; S30; S39; S57] calculate aggregate mode-share changes for different policy options, by modifying the variables of interest across the dataset, and rerunning the model.

4.6.1. Model selection - limitations

Two limitations are identified in relation to analysis of the final models in the studies in the review. One limitation is technical: studies not analysing uncertainty in performance estimates; and one is general: a limited understanding of how to use ML classifiers to inform policy decisions.

Q5a identifies that 69 out of the 73 studies do not analyse the expected distribution of the performance estimates. Each evaluation of model performance on a validation sample (whether through holdout validation or repeated cross-validation) is a random variable. If the distributions of the performance estimates are not accounted for, any apparent differences between different classifiers' performance estimates may be due to noise in this variable. Whilst several papers discuss the relative performance of classifiers for the mode choice prediction task, only one [S48] applies any formal test to investigate the statistical significance of differences between the classifiers. Additionally, as a discontinuous scoring metric (accuracy) is used and the number evaluations is low (10 folds of cross-validation) the direct distribution of the metric cannot be analysed, and instead non-parametric pairwise testing is used. This limitation represents an area for improvement.

Q5b identifies that there is currently a limited understanding in the literature of how to use ML classifiers to inform policy decisions. Few papers in the review attempt to extract behavioural indicators from ML classifiers, beyond the model's structural information. These are key outputs from RUMs that are used to inform policy making decisions. There is therefore a need to investigate further how these models can be used aside from prediction, for example to inform policy changes.

5. Conclusions

This paper conducts a systematic review of ML methodologies for modelling passenger mode choice. The review investigates five research questions covering classification techniques, datasets, performance estimation, model optimisation, and model selection.

A comprehensive search methodology across the three largest online publication databases is designed and used to identify 574 unique records. The record titles, abstracts, and publication details are screened for relevance, leaving 116 articles. The technical

Table 13
Limitations identified within systematic review.

No.	Classification stage	Description	Type
	Technical limitations		
TL1	Datasets	Studies not including any attributes of the mode-alternatives	Area for improvement
TL2	Datasets	Studies using input features which are dependent on output choice	Pitfall
TL3	Model validation	Studies using inappropriate validation schemes	Pitfall
TL4	Model validation	Studies using incorrect sampling methods for hierarchical data	Pitfall
TL5	Model validation	Studies not performing external validation	Area for improvement
TL6	Model validation	Studies using only discrete metrics	Pitfall
TL7	Model optimisation	Studies not performing any type of hyper-parameter optimisation	Area for improvement
TL8	Model optimisation	Studies not using rigorous hyper-parameter search schemes	Area for improvement
TL9	Model optimisation	Studies optimising hyper-parameters on test data	Pitfall
TL10	Model analysis	Studies not analysing uncertainty in performance estimates	Area for improvement
	General limitations		
GL1	Classification algorithms	Inconsistent representation of logit models in ML studies	
GL2	Datasets	Not describing the dataset and modelling process in sufficient detail	
GL3	Datasets	Lack of relevant, openly available datasets including mode-alternative attributes	
GL4	Datasets	Not considering sampling of the data from the population	
GL5	Model optimisation	Not presenting specific model hyper-parameters	
GL6	Model analysis	Limited understanding of how to use ML classifiers to inform policy decisions	

content of the full-text of these articles is assessed according to the eligibility criteria. In total, following the two screening processes, 70 full text peer-reviewed articles containing 73 primary studies are used for data extraction.

The studies are each reviewed in detail to extract 17 attributes covering the five research questions. Through this process, 16 limitations are identified: 10 technical limitations, and six general limitations. The limitations are summarised in Table 13. As shown in the Table 13, each technical limitation belongs to one of the classification stages out of classification techniques, datasets, performance estimation, model optimisation, and model selection.

Of the 10 limitations, five represent *pitfalls* in the methodologies which are likely to result in unreliable estimates of model performance and impact the results of an investigation, and five are identified as *areas for improvement* which are not strictly incorrect but could be addressed in order to improve the reliability of the results and/or predictive performance of the models.

A full summary of the technical limitations present in each study is given in Table 14. All studies have at least three technical limitations within their methodology, and only one study does not have any of the *pitfalls* [S49].

The prevalence of the limitations identified in this review highlights the need for a deeper understanding of the methodologies used for ML modelling of choice behaviour. Whilst experimental assessment of the implications of these limitations is left to further work, it is clear that several of the pitfalls violate the central *holdout validation* principle of ML classification. In particular, TL4 (*Studies using incorrect sampling methods for hierarchical data*), which is present in all studies which use unsampled trip-diary data, has serious *data leakage* implications, as the model is essentially validated on data it has observed during model training.

5.1. Recommendations and further work

As this paper shows, there is increasing research focus on ML as an alternative to RUMs for modelling passenger mode choice. This approach has the potential to provide valuable new insights into mode choice modelling research questions when used correctly. However, from the analysis in the systematic review, it is clear that the methodologies used are highly fragmented, and there needs to be further work to establish good standard methodological practice for the use of ML for choice modelling. In particular, almost all of the studies identified in the review show at least one of five methodological pitfalls identified, which will result in biased estimates of model performance. This review has not performed any quantitative analysis of the impacts of these pitfalls, or of the relative performance of the classifiers considered.

As identified by the general limitations, there is inconsistent representation of RUMs within papers that compare ML and RUMs. Furthermore, there is a limited understanding of how to use ML classifiers to inform policy decisions.

This leaves four key directions for further work: (i) establish a standardised methodology which can be used for both ML and random utility approaches which addresses the limitations raised in this review, (ii) use the methodology to investigate the impacts of the identified pitfalls on modelling results, (iii) use the methodology evaluate ML and RUM approaches to quantify fairly the trade-off in terms of predictive ability, and (iv) investigate how ML approaches can be used to inform policy changes and/or assist in specification of RUMs.

5.2. Limitations of systematic review

This section analyses the limitations of the review with respect to the recommended PRISMA guidelines (Moher et al., 2009). This review focuses on ML classification methodologies for modelling passenger mode choice. It therefore does not cover

Table 14
Summary of limitations within each study in systematic review.

No.	Paper	TL1	TL2	TL3	TL4	TL5	TL6	TL7	TL8	TL9	TL10	Sum
S1	Raju et al. (1996)				?	✓	✓		✓	?	✓	6
S2	Subba Rao et al. (1998)					✓	✓	✓	✓		✓	5
S3	Hensher and Ton (2000						?		✓	?	✓	4
S4	Cantarella and de Luca (2003)					✓			✓	✓	✓	4
S5	Van Middelkoop, Borgers, and Timmermans (2003	✓		✓	✓	✓	✓		✓	✓	✓	8
S6	Xie et al. (2003)	✓	✓		?	✓	✓		✓		✓	7
S7.1	Karlaftis (2004)		✓			✓	✓	✓	✓		✓	6
S7.2	–		✓	✓		✓	✓	✓	✓		✓	7
S7.3	–		✓			✓	✓	✓	✓		✓	6
S8.1	Cantarella and de Luca (2005)					✓			✓	✓	✓	4
S8.2	–					✓			✓	✓	✓	4
S9	Andrade et al. (2006)				✓	✓	✓	NA	NA		✓	4
S10	Shafahi and Nazari (2006				?	✓	✓	NA	NA		✓	4
S11	Edara, Teodorovi'c, and Baik (2007					✓	NA		✓	?	✓	4
S12	Errampalli et al. (2007)				?	✓	✓	NA	NA		✓	4
S13	Moons et al. (2007)					✓	✓	✓	✓		✓	5
S14	Wang and Namgung (2007)	✓			✓	✓	✓	NA	NA		✓	5
S15	Zhang and Xie (2008)				?	✓	✓		✓		✓	5
S16	Biagioni et al. (2009)		✓	✓	✓	✓	✓		✓	?	✓	8
S17	Chalumuri et al. (2009)				?	✓	✓		✓		✓	5
S18	Seetharaman et al. (2009)					✓	✓	NA	NA		✓	3
S19	Lu and Kawamura (2010)				✓	✓	✓	NA	NA		✓	4
S20	Zhao et al. (2010)					✓	✓		✓	?	✓	5
S21	Xian-Yu (2011)	✓			?	✓	✓		✓		✓	6
S22	Yin and Guan (2011)	?	?		?	✓	✓		✓	?	✓	8
S23	Zenina and Borisov (2011)	✓	✓	✓		✓	✓	✓	✓		✓	8
S24	Zhou and Lu (2011)	✓			?	✓	✓	✓	✓		✓	7
S25	Dell'Orco and Ottomanelli (2012)					✓	✓	NA	NA		✓	3
S26	Tang et al. (2012)	✓			?	✓	✓		✓	?		6
S27	Gao et al. (2013)	✓			?	✓	✓	✓	✓		✓	7
S28	Kumar et al. (2013)					✓	✓	NA	NA		✓	3
S29	Omrani et al. (2013)				✓	✓	✓		✓	?	✓	6
S30	Pulugurta et al. (2013)				✓	✓	✓	NA	NA		✓	4
S31	Ramanuj and Gundaliya (2013)	✓	✓		?	✓	✓		✓	?	✓	8
S32	Shukla et al. (2013)	✓	✓		✓	✓	✓	✓	✓		✓	8
S33	Cheng et al., 2014	✓			✓	✓	✓	NA	NA		✓	5
S34	Hossein et al. (2014)	✓			✓	✓	✓	✓	✓		✓	7
S35	Rasouli and Timmermans (2014)	✓			✓	✓	✓	✓	✓		✓	7
S36	Ermagun, Rashidi, and Lari (2015)					✓	✓	✓	✓		✓	5
S37	Gazder and Ratrout (2015)	✓	✓			✓	✓		✓	?	✓	7
S38	Jia et al. (2015)	✓			?	✓			✓	?	✓	6
S39	Kedia, Saw, and Katti (2015)	✓			?	✓	✓	NA	NA		✓	5
S40	Ma (2015)	✓	✓		✓	✓	✓	NA	NA		✓	6
S41	Omrani (2015)				✓	✓			✓	✓	✓	5

Table 14 (*continued*)

No.	Paper	TL1	TL2	TL3	TL4	TL5	TL6	TL7	TL8	TL9	TL10	Sum
S42	Papaiolannou and Martinez (2015)			✓	✓	✓	✓	NA	NA		✓	5
S43	Pitombo et al. (2015)	✓	✓			✓	✓	✓	✓		✓	7
S44	Tang et al. (2015)		✓		✓	✓	✓		✓		✓	6
S45	Li et al. (2016)	✓	✓	✓		✓	✓	✓	✓		✓	8
S46	Sekhar, Minal, and Madhu (2016)	?	?		?	✓		✓	✓		✓	7
S47	Semanjski et al. (2016)	✓				✓	✓		✓	?	✓	6
S48	Hagenauer and Helbich (2017)	✓			✓	✓	✓		✓			5
S49	Hussain et al. (2017)	✓				✓			✓		✓	4
S50	Juremalani (2017)	✓		✓		✓	✓	✓	✓		✓	7
S51	Lindner, Pitombo, and Cunha (2017)	✓				✓	✓	✓	✓		✓	6
S52	Ma et al. (2017)		✓			✓		NA	NA		✓	3
S53	Nam et al. (2017)					✓			✓	✓		3
S54	Assi, Nahiduzzaman, et al. (2018)	✓	✓			✓	✓		✓	?	✓	7
S55	Ding et al. (2018)	✓		✓	NA	✓	NA		✓	✓	✓	6
S56	Golshani et al. (2018)				✓	✓	✓		✓	✓	✓	6
S57	Lee et al. (2018)				✓	✓	✓		✓	?	✓	6
S58	Liang et al. (2018)	✓	✓	✓		✓	✓	✓	✓		✓	8
S59	Srivastava and Ravi Sekhar (2018)			✓		✓	✓		✓		✓	5
S60	Wang and Ross (2018)		✓		✓	✓	✓		✓	✓		6
S61	Zhu et al. (2018)				✓	✓		✓	✓		✓	5
S62	Assi and Shafiullah (2019)	✓	✓			✓	✓		✓	?	✓	7
S63	Chang et al. (2019)	✓	✓		✓	✓	✓		✓	?	✓	8
S64	Chapleau et al. (2019)		✓		✓		✓		✓	?	✓	6
S65	Cheng et al., 2019		✓		✓	✓	✓		✓		✓	6
S66	Minal et al. (2019)	?	?		?	✓	✓		✓	?	✓	8
S67	Pirra and Diana (2019) Wang and Zhao (2019)	✓			✓	✓	✓	✓	✓		✓	7
S68	Wang and Zhao (2019)				✓	✓	✓		✓		✓	5
S69	Yang and Ma (2019)		✓		✓	✓	✓	✓	✓		✓	7
S70	Zhou et al. (2019)	?	?	✓		✓	✓		✓	?	✓	8
Sum		34	26	13	40	71	61	21	59	28	69	

✓: Limitation present, ?: Unclear from text, NA: Not appicable.

contributions related to other ML techniques that have been used to investigate mode choice modelling, including clustering, reinforcement learning, and generative models. Furthermore, the review does not include ML applications for other choice modelling applications. However, the authors believe the findings of the review are relevant to other choice modelling domains.

Whilst a comprehensive and exhaustive search methodology covering the three largest online databases is used to identify relevant literature, there may have been relevant studies which are not included. In particular, the requirement for phrases related to mode choice in the title (used to pre-screen irrelevant articles) may have omitted studies which are of relevance to the review. Additionally, the review does not consider grey literature or unpublished material. However, in this new, research-led field, the authors are confident that the state-of-the-art techniques are well covered by the sample of studies assembled.

This review focuses purely on the methodologies used in each study and makes no attempt to draw conclusions on the findings reported by each paper. As such, no assessment is made of the quality of each paper, nor the publication bias of the field.

Whilst the procedure for the review is designed to be as objective as possible, the data extraction and discussion is carried out by the first author, under the guidance of the co-authors. This is according to available resources. All results and decisions have been double checked, but there may be remaining errors, which are the responsibility of the authors.

CRediT authorship contribution statement

Tim Hillel: Conceptualization, Investigation, Writing - original draft. **Michel Bierlaire:** Writing - review & editing, Supervision. **Mohammed Z.E.B. Elshafie:** Supervision. **Ying Jin:** Writing - review & editing, Supervision.

Declaration of competing interest

None.

Acknowledgements

References

Andrade, Katia, Uchida, Kenetsu, Kagaya, Seiichi, 2006. Development of transport mode choice model by using adaptive neuro-fuzzy inference System. Transport. Res. Rec. (1977), 8–16.

Assi, Khaled J., Nahiduzzaman, Kh Md, et al., 2018. Mode choice behavior of high school goers: evaluating logistic regression and MLP neural networks. Case Stud. Transp. Pol. 6 (2), 225–230.

Assi, Khaled J., Shafiullah, Md, et al., 2019. Travel-to-school mode choice modelling employing artificial intelligence techniques: a comparative study. Sustainability 11 (16), 4484.

Axhausen, Kay W., et al., 2002. Observing the rhythms of daily life: a six-week travel diary. Transportation 29, 96–124.

Barff, Richard, Mackay, David, Olshavsky, Richard W., 1982. A selective review of travel-mode choice models. J. Consum. Res. 8 (4), 370–380.

Ben-Akiva, Moshe, E., Lerman, Steven R., 1985. Discrete Choice Analysis: Theory and Application to Travel Demand. MIT Press, p. 424.

Bergstra, James, Bengio, Yoshua, 2012. Random search for hyper-parameter optimization. J. Mach. Learn. Res. 13 (Feb), 281–305.

Biagioni, James P., et al., 2009. Tour-based mode choice modeling: using an ensemble of conditional and unconditional data mining classifiers. In: Transportation Research Board 88th Annual Meeting, vol. 312. Transportation Research Board, Washington DC, USA, pp. 1–15.

Bierlaire, Michel, 2018. Biogeme examples - swissmetro (visited on 08/26/2018). http://biogeme.epfl.ch/examples_swissmetro.html.

Bierlaire, Michel, Lotan, Tsippy, Toint, Philippe, 1997. On the overspecification of multinomial and nested logit models due to alternative specific constants. Transport. Sci. 31 (4), 363–371.

Bierlaire, Michel, Kay, Axhausen, Abay, Georg, 2001. The acceptance of modal innovation: the case of swissmetro. In: Swiss Transport Research Conference.

Breiman, Leo, 2017. Classification and Regression Trees. Routledge.

Bureau of Transportation Statistics, 2018. The 1995 American travel survey (ATS) - household trip characteristics (visited on 08/26/2018). https://catalog.data.gov/dataset/the-1995-american-travel-survey-ats-household-trips.

Cantarella, Giulio Erberto, de Luca, Stefano, 2003. Modeling transportation mode choice through artificial neural networks. In: Fourth International Symposium on Uncertainty Modeling and Analysis, 2003. ISUMA 2003, College Park, MD, USA, pp. 84–90. IEEE.

Cantarella, Giulio Erberto, de Luca, Stefano, 2005. Multilayer feedforward networks for transportation mode choice analysis: an analysis and a comparison with random utility models. In: Transportation Research Part C: Emerging Technologies. Handling Uncertainty in the Analysis of Traffic and Transportation Systems. Bari, Italy, pp. 121–155. June 10-13 2002) 13.2.

Chalumuri, Ravi Sekhar, et al., 2009. Applications of neural networks in mode choice modelling for second order metropolitan cities of India. In: Proceedings of the Eastern Asia Society for Transportation Studies, vol. 7. Eastern Asia Society for Transportation Studies, Surabaya, Indonesia, pp. 1–16.

Chang, Ximing, et al., 2019. Travel mode choice: a data fusion model using machine learning methods and evidence from travel diary survey data. Transportmetrica A - Transp. Sci. 15, 1587–1612, 2.

Chapleau, R., Gaudette, P., Spurr, T., 2019. Application of machine learning to two large-sample household travel surveys: a characterization of travel modes. Transport. Res. Rec.: J. Transp. Res. Board 2673 (4), 173–183.

Cheng, Long, Chen, Xuewu, Ming, Wei, et al., 2014. Modeling mode choice behavior incorporating household and individual sociodemographics and travel attributes based on Rough sets theory. Comput. Intell. Neurosci. (2014), 1–9.

Cheng, Long, Chen, Xuewu, Jonas De Vos, et al., 2019. Applying a random forest method approach to model travel mode choice behavior. Travel Behav. Soc. 14, 1–10.

Chicago Metropolitan Agency for Planning, 2018a. CATS household travel survey, 1990 (visited on 08/26/2018). https://datahub.cmap.Illinois.gov/dataset/travel-survey-1990.

Chicago Metropolitan Agency for Planning, 2018b. Travel tracker survey, 2007 - 2008: public data. visited on 08/26/2018. https://datahub.cmap.illinois.gov/dataset/traveltracker0708.

City of Chicago, 2020. Taxi trips. visited on 01/27/2020. https://data.cityofchicago.org/Transportation/Taxi-Trips/wrvz-psew.

Cortes, Corinna, Vapnik, Vladimir, 1995. Support-vector networks. Mach. Learn. 20, 273–297, 3.

Croissant, Yves, 2016. Ecdat: data sets for econometrics (visited on 08/26/2018). https://CRAN.R-project.org/package=Ecdat.

Delaware Valley Regional Planning Commission, 2018. 2012 Household Travel Survey. URL. https://www.dvrpc.org/transportation/Modeling/Data/ (visited on 08/26/2018).

Dell'Orco, Mauro, Ottomanelli, Michele, 2012. Simulation of users decision in transport mode choice using neuro-fuzzy approach. In: International Conference on Computational Science and its Applications (ICCSA 2012). Salvador de Bahia, Brazil, pp. 44–53. Springer.

Ding, Chuan, Cao, Xinyu, Wang, Yunpeng, 2018. Synergistic effects of the Built environment and commuting programs on commute mode choice. Transport. Res. Pol. Pract. vol. 118, 104–118.

Divvy Bikes, 2020. Divvy System data (visited on 01/27/2020). https://www.divvybikes.com/system-data.

Edara, Praveen Kumar, Teodorović, Dusan, Baik, Hojong, 2007. Using neural networks to model intercity mode choice. In: Smart Systems Engineering: Computational Intelligence in Architect- Ing Complex Engineering Systems, vol. 17. ASME Press, St Louis, Missouri, USA, pp. 143–148. Artificial Neural Networks in Engineering Conference (ANNIE 2007).

Ermagun, Alireza, Rashidi, Taha Hossein, Ansari Lari, Zahra, 2015. Mode choice for school trips: long-term planning and impact of modal specification on policy assessments. Transport. Res. Rec. vol. 2513, 97–105.

Errampalli, Madhu, Okushima, Masashi, Akiyama, Takamasa, 2007. Combined fuzzy logic based mode choice and microscopic simulation model for transport policy evaluation. In: 11th World Conference on Transport Research. Transportation Research Board, Berkley CA, USA.

Gao, Jian, et al., 2013. Impact of transit network layout on resident mode choice. Math. Probl Eng. (2013), 1–8.

Gazder, Uneb, Ratrout, Nedal T., 2015. A new logit-artificial neural network ensemble for mode choice modeling: a case study for border transport. J. Adv. Transport. 49 (8), 855–866.

Gneiting, Tilmann, Raftery, Adrian E., 2007. Strictly proper scoring rules, prediction, and estimation. J. Am. Stat. Assoc. 102 (477), 359–378.

Golshani, Nima, et al., 2018. Modeling travel mode and timing decisions: comparison of artificial neural networks and copula-based joint model. In: Travel Behaviour and Society, vol. 10, pp. 21–32.

Greene, William H., 2011. Econometric Analysis. Pearson Higher, p. 1230.

Hagenauer, Julian, Helbich, Marco, 2017. A comparative study of machine learning classifiers for modeling travel mode choice. Expert Syst. Appl. vol. 78, 273–282.

Hartigan, J.A., Wong, M.A., 1979. Algorithm as 136: a K-means clustering algorithm. J. R. Statist. Soc. Ser. C (Appl. Statist.) 28, 100–108, 1.

Hastie, Trevor, Friedman, Jerome, Tibshirani, Robert, 2008. The Elements of Statistical Learning, second ed. Springer Series in Statistics, New York: Springer, p. 745.

Hensher, David A., Johnson, Lester W., 1983. Alternative modelling procedures in studies of travel mode choice: a review and appraisal. Transport. Plann. Technol. 8 (3), 203–216.

Hensher, David A., Ton, Tu T., 2000. A comparison of the predictive potential of artificial neural networks and nested logit models for commuter mode choice. Transport. Res. E Logist. Transport. Rev. 36 (3), 155–172.

Federal Highway Administration, 2018. National household travel survey (visited on 08/26/2018). https://nhts.ornl.gov/.

Hillel, Tim, Elshafie, Mohammed ZEB., Jin, Ying, 2018. Recreating passenger mode choice-sets for transport simulation: a case study of london, UK. Proc. Inst. Civ. Eng. Smart Infrastruct. Construct. 171 (1), 29–42.

Hoos, Holger, et al., 2014. An efficient approach for assessing hyperparameter importance. In: International Conference on Machine Learning, pp. 754–762.

Rashidi, Hossein, Taha, Hasegawa, Hironobu, 2014. An innovative simultaneous System of disaggregate models for trip generation, mode, and destination choice. In: Transportation Research Board 93rd Annual Meeting. Transportation Research Board, Washington DC, USA.

Hussain, H.D., et al., 2017. Analysis of transportation mode choice using a comparison of artificial neural network and multinomial logit models. ARPN J. Eng. Appl. Sci. 12 (5), 1483–1493.

Jia, Hongfei, Cao, Xiongjiu, Yang, Kaihua, 2015. Residents' travel mode choice model. Traffic Eng. Contr. 56 (1), 169–174.

Jing, Peng, et al., 2018. Travel mode and travel route choice behavior based on random regret minimization: a systematic review. Sustainability 10 (4), 1185.

Juremalani, Jayesh, 2017. Comparison of different mode choice models for work trips using data mining process. Indian J. Sci. Technol. 10 (17), 1–3.

Karlaftis, Matthew G., 2004. Predicting mode choice through multivariate recursive partitioning. J. Transport. Eng. 130 (2), 245–250.

Kedia, Ashu Shivkumar, Saw, Krishna Bhuneshwar, Katti, Bhimaji Krishnaji, 2015. Fuzzy logic approach in mode choice modelling for education trips: a case study of Indian metropolitan city. Transport 30.3, 286–293.

Kitchenham, Barbara, Charters, Stuart, 2007. Guidelines for Performing Systematic Literature Reviews in Software Engineering. EBSE. EBSE Technical Report 2007-01.

Kruger, J., 1991. Review of research on urban area mode choice modelling. In: 13th Conference of Australian Institutes of Transport Research (CAITR). Transportation Research Board, Brisbane, Queensland.

Kumar, Mukesh, Sarkar, Pradip, Madhu, Errampalli, 2013. Development of fuzzy logic based mode choice model considering various public transport policy options. Int. J. Traffic Transport. Eng. 3 (4), 408–425.

Lee, Dongwoo, Derrible, Sybil, Pereira, Francisco Camara, 2018. Comparison of four types of artificial neural networks and a multinomial logit model for travel mode choice modeling. In: Transportation Research Board 97th Annual Meeting. Transportation Research Board, Washington DC, USA.

Li, Juan, et al., 2016. Cluster-based logistic regression model for holiday travel mode choice. In: Procedía Engineering. 6th International Conference on Green Intelligent Transportation System and Safety (GITSS 2015), vol. 137. Elsevier, Beijing, China, pp. 729–737.

Liang, LeiLei, et al., 2018. Travel mode choice analysis based on household mobility survey data in milan: comparison of the multinomial logit model and random forest approach. In: Transportation Research Board 97th Annual Meeting. Transportation Research Board, Washington DC, USA.

Lindner, Anabele, Souza Pitombo, Cira, Cunha, André Luiz, 2017. Estimating motorized travel mode choice using classifiers: an application for high-dimensional multicollinear data. In: Travel Behaviour and Society, vol. 6, pp. 100–109.

Lu, Yandan, Kawamura, Kazuya, 2010. Data-mining approach to work trip mode choice analysis in Chicago, Illinois, area. In: Transportation Research Record, vol. 2156, pp. 73–80.

Luxembourg Institute of Socio-Economic Research, 2018. Socio-economic panel ofLiewen zu lëtze- buerg III (PSELL3) (visited on 08/26/2018). http://dataservice.liser.lu/en_US/dataservice/db=23.

Ma, Tai-Yu, 2015. Bayesian networks for multimodal mode choice behavior modelling: a case study for the cross border workers of Luxembourg. In: Transportation Research Procedia. 18th Euro Working Group on Transportation, EWGT 2015, 14-16 July 2015, vol. 10. Elsevier, Delft, The Netherlands, pp. 870–880. Delft, The Netherlands.

Ma, Tai-Yu, Joseph, Y., Chow, J., Jia, Xu, 2017. Causal structure learning for travel mode choice using structural restrictions and model averaging algorithm. Transportmetrica: Transport. Sci. 13 (4), 299–325.

McFadden, Daniel, 1981. Econometric models of probabilistic choice. In: Manski, Charles F., McFadden, Daniel (Eds.), Structural Analysis of Discrete Data with Econometric Applications. MIT Press, pp. 198–272.

Meixell, Mary J., Norbis, Mario, 2008. A review of the transportation mode choice and carrier selection literature. Int. J. Logist. Manag. 19 (2), 183–211.

Metropolitan Transportation Commission, 2018a. 1990 Bay area travel surveys (visited on 08/26/2018). http://www.surveyarchive.org/.

Metropolitan Transportation Commission, 2018b. San Francisco Bay area travel survey 2000 (visited on 08/26/2018). http://www.surveyarchive.org/.

Minal, Ravi Sekhar, Ch, 2014. Mode choice analysis: the data, the models and future ahead. Int. J. Traffic Transport. Eng. 4 (3), 269–285.

Minal, S., Sekhar, Ch Ravi, Madhu, Errampilli, 2019. Development of neuro-fuzzy-based multimodal mode choice model for commuter in Delhi. IET Intell. Transp. Syst. 13 (2), 243–251.

Moher, David, et al., 2009. Preferred reporting Items for systematic reviews and meta-analyses: the PRISMA statement. PLoS Med. 6 (7), 6.

Moons, Elke, Wets, Geert, Aerts, Marc, 2007. Nonlinear models for determining mode choice. In: Prog. Artif. Intell. 13th Portuguese Conference on Aritficial Intelligence (EPIA 2007). Lecture Notes in Computer Science. Springer, Guimaraes, Portugal, pp. 183–194.

Nam, Daisik, et al., 2017. A model based on deep learning for predicting travel mode choice. In: Transportation Research Board 96th Annual Meeting. Transportation Research Board, Washington DC, USA, pp. 8–12.

Nerhagen, L., 2000. Mode choice behaviour, travel mode choice models and value of time esti- amtion. A literature review. In: CTEK Working Paper.

Niculescu-Mizil, Alexandru, Caruana, Rich, 2005. Predicting good probabilities with supervised learning. In: Proceedings of the 22nd International Conference on Machine Learning. ACM, pp. 625–632.

Omrani, Hichem, 2015. Predicting travel mode of individuals by machine learning. In: Transportation Research Procedia, vol. 10. Elsevier, Delft, The Netherlands, pp. 840–849.

Omrani, Hichem, et al., 2013. Prediction of individual travel mode with evidential neural network model. Transport. Res. Rec. 2399, 1–8.

Papaioannou, Dimitrios, Martinez, Luis Miguel, 2015. The role of accessibility and connectivity in mode choice. A structural equation modeling approach. In: Transportation Research Procedia.Vol. 10. 18th Euro Working Group on Transportation, EWGT 2015. Elsevier, Delft, The Netherlands, pp. 831–839.

Pirra, Miriam, Diana, Marco, 2019. A study of tour-based mode choice based on a support vector machine classifier. Transport. Plann. Technol. 42 (1), 23–36.

Pitombo, Souza, Cira, et al., 2015. A two-step method for mode choice estimation with socioeconomic and spatial information. Spatial Statist. 11, 45–64.

Platt, John C., 1999. Probabilistic outputs for support vector machines and comparisons to regularized likelihood methods. In: Advances in Large Margin Classifiers. MIT Press, pp. 61–74.

Pulugurta, Sarada, Arun, Ashutosh, Errampalli, Madhu, 2013. Use of artificial intelligence for mode choice analysis and comparison with traditional multinomial logit model. In: Procedia - Social and Behavioral Sciences. 2 nd Conference of Transportation Research Group of India (CTRG 2013), vol. 104. Elsevier, Agra, India, pp. 583–592.

Raju, K.A., Sikdar, P.K., Dhingra, S.L., 1996. Modelling mode choice by means of an artificial neural network. Environ. Plann. Plann. Des. 23 (6), 677–683.

Ramanuj, P.S., Gundaliya, P.J., 2013. Disaggregated modeling of mode choice by ANN-a case study of ahmedabad city in Gujarat state. J. Indian Roads Congress 74 (1), 3–12.

Rasouli, Soora, Timmermans, Harry J.P., 2014. Using ensembles of decision trees to predict transport mode choice decisions: effects on predictive success and uncertainty estimates. Eur. J. Transport Infrastruct. Res. 14 (4), 412–424.

Ratrout, Nedal T., Gazder, Uneb, Al-Madani, Hashim M.N., 2014. A review of mode choice modelling techniques for intra-city and border transport. World Rev. Intermodal Transp. Res. 5 (1), 39–58.

Saeb, Sohrab, et al., 2017. The need to approximate the use-case in clinical machine learning. GigaScience 6 (5), 1–9.

Seetharaman, Padma, et al., 2009. Comparative evaluation of mode choice modelling by logit and fuzzy logic. In: Proceedings of the Eastern Asia Society for Transportation Studies, vol. 7. Eastern Asia Society for Transportation Studies, Surabaya, Indonesia, pp. 1–16.

Sekhar, Ch, Ravi, Minal, Madhu, E., 2016. Mode choice analysis using random forrest decision trees. In: Transportation Research Procedia, vol. 17. Elsevier, Mumbai, India, pp. 644–652, 11th Transportation Planning and Implementation Methodologies for Developing Countries, TPMDC 2014.

Semanjski, Ivana, Lopez, Angel, Gautama, Sidharta, 2016. Forecasting transport mode use with support vector machines based approach. Trans. Marit. Sci. 5 (2), 111–120.

Shafahi, Yusof, Nazari, Sobhaan, 2006. Disaggregate mode choice analysis for work trips using genetic-fuzzy and neuro-fuzzy systems. In: Proceedings of the 10th IASTED International Conference on Artificial Intelligence and Soft Computing (ASC 2006). ACTA Press, Spain, pp. 250–255. Palma de Mallorca.

Shmueli, Galit, 2010. To explain or to predict? Stat. Sci. 25 (3), 289–310.

Shukla, Nagesh, et al., 2013. Data-driven modeling and analysis of household travel mode choice. In: 20th International Congress on Modelling and Simulation (MODSIM 2013). The Modelling and Simulation Society of Australia and New Zealand Inc., Adelaide, Australia, pp. 92–98.

Srivastava, Nitish, Hinton, Geoffrey, et al., 2014. Dropout: a simple way to prevent neural networks from overfitting. J. Mach. Learn. Res. 15 (1), 1929–1958.

Srivastava, Minal, Ravi Sekhar, Chalumuri, 2018. Web survey data and commuter mode choice analysis using artificial neural network. Int. J. Traffic Transport. Eng. 8 (3), 359–371.

Subba Rao, P.V., et al., 1998. Another insight into artificial neural networks through behavioural analysis of access mode choice. Comput. Environ. Urban Syst. 22 (5), 485–496.

Svozil, Daniel, Kvasnicka, Vladimir, Pospichal, Jiri, 1997. Introduction to multi-layer feed-forward neural networks. Chemometr. Intell. Lab. Syst. 39 (1), 43–62.

Tang, Dounan, Yang, Min, Zhang, Mei Hui, 2012. Travel mode choice modeling: a comparison of bayesian networks and neural networks. Appl. Mech. Mater. 209–211, 717–723.

Tang, Liang, Xiong, Chenfeng, Zhang, Lei, 2015. Decision tree method for modeling travel mode switching in a dynamic behavioral process. Transport. Plann. Technol. 38 (8), 833–850.

Train, K.E., 2009. Discrete Choice Methods with Simulation. Cambridge university press.

Transport for, N.S.W., 2018. Household Travel Survey 2005/06-2016/17 (visited on 08/26/2018). https://opendata.transport.nsw.gov.au/dataset/household-travel-survey-200506-%E2%80%93-201617.

Transport for Victoria, 2018. VISTA Data and Publications (visited on 08/26/2018). https://transport.vic.gov.au/data-and-research/vista/vista-data-and-publications/.

Van Middelkoop, Manon, Borgers, Aloys, Timmermans, Harry, 2003. Inducing heuristic principles of tourist choice of travel mode: a rule-based approach. J. Trav. Res. 42 (1), 75–83.

Varma, Sudhir, Simon, Richard, 2006. Bias in error estimation when using cross-validation for model selection. BMC Bioinf. 8.

Wang, Weijie, Moon, Namgung, 2007. Knowledge discovery from the data of long distance travel mode choices based on Rough set theory. In: International Journal of Multimedia and Ubiquitous Engineering, vol. 2, pp. 81–90, 3.

Wang, Fangru, Ross, Catherine L., 2018. Machine learning travel mode choices: comparing the performance of an extreme gradient boosting model with a multinomial logit model. Transport. Res. Record vol. 2672 (47), 35–45.

Wang, Shenhao, Zhao, Jinhua, 2019. An empirical study of using deep neural network to analyze travel mode choice with interpretable economic information. In: Transportation Research Board 98th Annual Meeting. Transportation Research Board, Washington DC, USA.

Wu, Ting-Fan, Lin, Chih-Jen, Ruby, C., Weng, 2004. Probability estimates for multi-class classification by pairwise coupling. J. Mach. Learn. Res. 5 (Aug), 975–1005.

Xian-Yu, Jian-Chuan, 2011. Travel mode choice analysis using support vector machines. In: 11th International Conference of Chinese Transportation Professionals (ICCTP 2011). American Society of Civil Engineers, Nanjing, China, pp. 360–371.

Xie, Chi, Lu, Jinyang, Parkany, Emily, 2003. Work travel mode choice modeling with data mining: decision trees and neural networks. Transp. Res. Record 1854, 50–61.

Yang, Jie, Ma, Jun, 2019. Compressive sensing-enhanced feature selection and its application in travel mode choice prediction. In: Applied Soft Computing, vol. 75, pp. 537–547.

Yin, Huanhuan, Guan, Hongzhi, 2011. Traffic mode choice model based on BP neural network. In: Proceedings 2011 International Conference on Transportation, Mechanical, and Electrical Engineering (TMEE). IEEE, Changchun, China, pp. 1441–1444.

Zenina, Nadezda, Borisov, Arkady, 2011. Transportation mode choice analysis based on classification methods. In: Scientific Journal of Riga Technical University. Computer Sciences, vol. 45, pp. 49–53, 1.

Zhang, Yunlong, Xie, Yuanchang, 2008. Travel mode choice modeling with support vector machines. Transp. Res. Record 2076, 141–150.

Zhao, Dan, et al., 2010. Travel mode choice modeling based on improved probabilistic neural network. In: Traffic and Transportation Studies 2010 (ICTTS 2010), vol. 383. Kunming, China: ASCE, pp. 685–695.

Zhou, Miaomiao, Lu, Jian, 2011. Research on prediction of traffic mode choice of urban residents. In: 11th International Conference of Chinese Transportation Professionals (ICCTP 2011). American Society of Civil Engineers, Nanjing, China, pp. 449–460.

Zhou, Xiaolu, Wang, Mingshu, Li, Dongying, 2019. Bike-sharing or taxi? Modeling the choices of travel mode in Chicago using machine learning. In: Journal of Transport Geography, vol. 79. UNSP 102479.

Zhu, Zheng, et al., 2018. A mixed bayesian network for two-dimensional decision modeling of departure time and mode choice. Transportation 45, 1499–1522.